

Problem Statement

Business Context

Business communities in the United States are facing high demand for human resources, but one of the constant challenges is identifying and attracting the right talent, which is perhaps the most important element in remaining competitive. Companies in the United States look for hard-working, talented, and qualified individuals both locally as well as abroad.

The Immigration and Nationality Act (INA) of the US permits foreign workers to come to the United States to work on either a temporary or permanent basis. The act also protects US workers against adverse impacts on their wages or working conditions by ensuring US employers' compliance with statutory requirements when they hire foreign workers to fill workforce shortages. The immigration programs are administered by the Office of Foreign Labor Certification (OFLC).

OFLC processes job certification applications for employers seeking to bring foreign workers into the United States and grants certifications in those cases where employers can demonstrate that there are not sufficient US workers available to perform the work at wages that meet or exceed the wage paid for the occupation in the area of intended employment.

Objective

In FY 2016, the OFLC processed 775,979 employer applications for 1,699,957 positions for temporary and permanent labor certifications. This was a nine percent increase in the overall number of processed applications from the previous year. The process of reviewing every case is becoming a tedious task as the number of applicants is increasing every year.

The increasing number of applicants every year calls for a Machine Learning based solution that can help in shortlisting the candidates having higher chances of VISA approval. OFLC has hired the firm EasyVisa for data-driven solutions. You as a data scientist at EasyVisa have to analyze the data provided and, with the help of a classification model:

- Facilitate the process of visa approvals.
- Recommend a suitable profile for the applicants for whom the visa should be certified or denied based on the drivers that significantly influence the case status.

Data Description

The data contains the different attributes of employee and the employer. The detailed data dictionary is given below.

- `case_id`: ID of each visa application
- `continent`: Information of continent the employee
- `education_of_employee`: Information of education of the employee
- `has_job_experience`: Does the employee has any job experience? Y= Yes; N = No
- `requires_job_training`: Does the employee require any job training? Y = Yes; N = No
- `no_of_employees`: Number of employees in the employer's company
- `yr_of_estab`: Year in which the employer's company was established
- `region_of_employment`: Information of foreign worker's intended region of employment in the US.
- `prevailing_wage`: Average wage paid to similarly employed workers in a specific occupation in the area of intended employment. The purpose of the prevailing wage is to ensure that the foreign worker is not underpaid compared to other workers offering the same or similar service in the same area of employment.
- `unit_of_wage`: Unit of prevailing wage. Values include Hourly, Weekly, Monthly, and Yearly.
- `full_time_position`: Is the position of work full-time? Y = Full Time Position; N = Part Time Position
- `case_status`: Flag indicating if the Visa was certified or denied

Installing and Importing the necessary libraries

```
In [1]: # Installing the Libraries with the specified version.  
# !pip install numpy==1.25.2 pandas==1.5.3 scikit-learn==1.5.2 matplotlib==3.  
7.1 seaborn==0.13.1 xgboost==2.0.3 -q --user
```

Note: After running the above cell, kindly restart the notebook kernel and run all cells sequentially from the below.

```
In [2]: # Libraries to help with reading and manipulating data
import pandas as pd
import numpy as np

# Libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# To tune model, get different metric scores, and split data
from sklearn.metrics import (
    f1_score,
    accuracy_score,
    recall_score,
    precision_score,
    confusion_matrix,
    roc_auc_score,
    ConfusionMatrixDisplay,
)
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score

# To be used for data scaling and one hot encoding
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder

# To impute missing values
from sklearn.impute import SimpleImputer
from sklearn import metrics

# To oversample and undersample data
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler

# To do hyperparameter tuning
from sklearn.model_selection import RandomizedSearchCV

# To define maximum number of columns to be displayed in a dataframe
pd.set_option("display.max_columns", None)

# To suppress scientific notations for a dataframe
pd.set_option("display.float_format", lambda x: "%.3f" % x)

# To help with model building
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (
    AdaBoostClassifier,
    GradientBoostingClassifier,
    RandomForestClassifier,
    BaggingClassifier,
)
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression

# To suppress scientific notations
pd.set_option("display.float_format", lambda x: "%.3f" % x)

# To suppress warnings
```

```
import warnings

warnings.filterwarnings("ignore")
```

Import Dataset

```
In [3]: from google.colab import drive
drive.mount('/content/drive')
file = r'/content/drive/MyDrive/AI Class/Projects/Project 3/EasyVisa.csv'
EasyVisaCampgn = pd.read_csv(file)
```

Mounted at /content/drive

Overview of the Dataset

View the first and last 5 rows of the dataset

```
In [4]: data = EasyVisaCampgn.copy()
```

```
In [5]: data.head()
```

Out[5]:

	case_id	continent	education_of_employee	has_job_experience	requires_job_training	no_of_e
0	EZYV01	Asia	High School	N	N	
1	EZYV02	Asia	Master's	Y	N	
2	EZYV03	Asia	Bachelor's	N	Y	
3	EZYV04	Asia	Bachelor's	N	N	
4	EZYV05	Africa	Master's	Y	N	

```
In [6]: data.tail()
```

Out[6]:

	case_id	continent	education_of_employee	has_job_experience	requires_job_training	
25475	EZYV25476	Asia	Bachelor's	Y	Y	
25476	EZYV25477	Asia	High School	Y	N	
25477	EZYV25478	Asia	Master's	Y	N	
25478	EZYV25479	Asia	Master's	Y	Y	
25479	EZYV25480	Asia	Bachelor's	Y	N	

Understand the shape of the dataset

```
In [7]: data.shape
```

```
Out[7]: (25480, 12)
```

Check the data types of the columns for the dataset

```
In [8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 25480 entries, 0 to 25479  
Data columns (total 12 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   case_id                              25480 non-null  object  
1   continent                             25480 non-null  object  
2   education_of_employee                 25480 non-null  object  
3   has_job_experience                     25480 non-null  object  
4   requires_job_training                 25480 non-null  object  
5   no_of_employees                       25480 non-null  int64  
6   yr_of_estab                           25480 non-null  int64  
7   region_of_employment                 25480 non-null  object  
8   prevailing_wage                       25480 non-null  float64  
9   unit_of_wage                          25480 non-null  object  
10  full_time_position                    25480 non-null  object  
11  case_status                           25480 non-null  object  
dtypes: float64(1), int64(2), object(9)  
memory usage: 2.3+ MB
```

```
In [9]: object_cols = data.select_dtypes(include='object').columns

for col in object_cols:
    data[col] = data[col].astype('category')

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25480 entries, 0 to 25479
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   case_id                               25480 non-null  category
1   continent                             25480 non-null  category
2   education_of_employee                 25480 non-null  category
3   has_job_experience                    25480 non-null  category
4   requires_job_training                 25480 non-null  category
5   no_of_employees                      25480 non-null  int64
6   yr_of_estab                          25480 non-null  int64
7   region_of_employment                 25480 non-null  category
8   prevailing_wage                      25480 non-null  float64
9   unit_of_wage                         25480 non-null  category
10  full_time_position                   25480 non-null  category
11  case_status                          25480 non-null  category
dtypes: category(9), float64(1), int64(2)
memory usage: 2.0 MB
```

```
In [10]: data.duplicated().sum()
```

```
Out[10]: np.int64(0)
```

```
In [11]: data.isna().sum()
```

```
Out[11]:
```

	0
case_id	0
continent	0
education_of_employee	0
has_job_experience	0
requires_job_training	0
no_of_employees	0
yr_of_estab	0
region_of_employment	0
prevailing_wage	0
unit_of_wage	0
full_time_position	0
case_status	0

dtype: int64

```
In [12]: data.isnull().sum()
```

```
Out[12]:
```

	0
case_id	0
continent	0
education_of_employee	0
has_job_experience	0
requires_job_training	0
no_of_employees	0
yr_of_estab	0
region_of_employment	0
prevailing_wage	0
unit_of_wage	0
full_time_position	0
case_status	0

dtype: int64

```
In [13]: data.nunique()
```

```
Out[13]:
```

	0
case_id	25480
continent	6
education_of_employee	4
has_job_experience	2
requires_job_training	2
no_of_employees	7105
yr_of_estab	199
region_of_employment	5
prevailing_wage	25454
unit_of_wage	4
full_time_position	2
case_status	2

dtype: int64

```
In [14]: round(data.isnull().sum() / data.isnull().count() * 100, 2)
```

```
Out[14]:
```

	0
case_id	0.000
continent	0.000
education_of_employee	0.000
has_job_experience	0.000
requires_job_training	0.000
no_of_employees	0.000
yr_of_estab	0.000
region_of_employment	0.000
prevailing_wage	0.000
unit_of_wage	0.000
full_time_position	0.000
case_status	0.000

dtype: float64

Exploratory Data Analysis (EDA)

Let’s check the statistical summary of the data

In [15]: data.describe().T

Out[15]:

	count	mean	std	min	25%	50%	75%	
no_of_employees	25480.000	5667.043	22877.929	-26.000	1022.000	2109.000	3504.000	€
yr_of_estab	25480.000	1979.410	42.367	1800.000	1976.000	1997.000	2005.000	
prevailing_wage	25480.000	74455.815	52815.942	2.137	34015.480	70308.210	107735.513	\$

In [16]: data.describe(include=['category']).T

Out[16]:

	count	unique	top	freq
case_id	25480	25480	EZYV9999	1
continent	25480	6	Asia	16861
education_of_employee	25480	4	Bachelor's	10234
has_job_experience	25480	2	Y	14802
requires_job_training	25480	2	N	22525
region_of_employment	25480	5	Northeast	7195
unit_of_wage	25480	4	Year	22962
full_time_position	25480	2	Y	22773
case_status	25480	2	Certified	17018

Fixing the negative values in number of employees columns

```
In [17]: negative_rows = data[data['no_of_employees'] < 0]
         print(negative_rows)

         #printing all rows with employees < 0
```

	case_id	continent	education_of_employee	has_job_experience	\
245	EZYV246	Europe	Master's	N	
378	EZYV379	Asia	Bachelor's	N	
832	EZYV833	South America	Master's	Y	
2918	EZYV2919	Asia	Master's	Y	
6439	EZYV6440	Asia	Bachelor's	N	
6634	EZYV6635	Asia	Bachelor's	Y	
7224	EZYV7225	Europe	Doctorate	N	
7281	EZYV7282	Asia	High School	N	
7318	EZYV7319	Asia	Bachelor's	Y	
7761	EZYV7762	Asia	Master's	N	
9872	EZYV9873	Europe	Master's	Y	
11493	EZYV11494	Asia	High School	Y	
13471	EZYV13472	North America	Master's	N	
14022	EZYV14023	Asia	Bachelor's	N	
14146	EZYV14147	Asia	Bachelor's	N	
14726	EZYV14727	Asia	Master's	N	
15600	EZYV15601	Asia	Bachelor's	N	
15859	EZYV15860	Asia	High School	N	
16157	EZYV16158	Asia	Master's	Y	
16883	EZYV16884	North America	Bachelor's	Y	
17006	EZYV17007	Asia	Doctorate	Y	
17655	EZYV17656	North America	Bachelor's	Y	
17844	EZYV17845	Asia	Bachelor's	N	
17983	EZYV17984	Asia	Bachelor's	N	
20815	EZYV20816	Asia	Bachelor's	N	
20984	EZYV20985	Europe	Doctorate	Y	
21255	EZYV21256	North America	High School	N	
21760	EZYV21761	Asia	Bachelor's	Y	
21944	EZYV21945	Africa	Master's	Y	
22084	EZYV22085	North America	Bachelor's	Y	
22388	EZYV22389	Asia	Master's	Y	
23186	EZYV23187	Asia	Master's	N	
23476	EZYV23477	Europe	Master's	Y	

	requires_job_training	no_of_employees	yr_of_estab	\
245	N	-25	1980	
378	Y	-11	2011	
832	N	-17	2002	
2918	N	-26	2005	
6439	N	-14	2013	
6634	N	-26	1923	
7224	N	-25	1998	
7281	N	-14	2000	
7318	Y	-26	2006	
7761	N	-11	2009	
9872	N	-26	1996	
11493	N	-14	1999	
13471	N	-17	2003	
14022	Y	-11	1946	
14146	Y	-26	1954	
14726	N	-11	2000	
15600	N	-14	2014	
15859	N	-11	1969	
16157	N	-11	1994	
16883	N	-26	1968	
17006	N	-11	1984	

17655	N	-17	2007
17844	N	-14	2012
17983	N	-26	2004
20815	Y	-17	1990
20984	N	-14	1989
21255	N	-25	1987
21760	N	-25	2000
21944	N	-25	1977
22084	N	-14	1980
22388	N	-14	1986
23186	Y	-11	2007
23476	N	-11	2000

	region_of_employment	prevailing_wage	unit_of_wage	full_time_position
245	Northeast	39452.990	Year	Y
378	Northeast	32506.140	Year	Y
832	South	129701.940	Year	Y
2918	Midwest	112799.460	Year	Y
6439	South	103.970	Hour	Y
6634	West	5247.320	Year	Y
7224	Midwest	141435.950	Year	Y
7281	Midwest	58488.500	Year	Y
7318	South	115005.610	Year	Y
7761	Midwest	38457.510	Year	Y
9872	South	37397.050	Year	Y
11493	South	27599.350	Year	Y
13471	Northeast	257.241	Hour	Y
14022	Northeast	108403.560	Year	Y
14146	West	81982.270	Year	Y
14726	Midwest	167851.800	Year	Y
15600	South	24641.610	Year	Y
15859	South	44640.600	Year	Y
16157	South	62681.250	Year	Y
16883	Northeast	168.156	Hour	Y
17006	West	25753.510	Year	Y
17655	Northeast	129753.180	Year	Y
17844	West	29325.850	Year	Y
17983	South	84359.980	Year	Y
20815	West	91897.570	Year	Y
20984	Midwest	37012.800	Year	Y
21255	South	99405.470	Year	N
21760	West	100463.580	Year	Y
21944	Midwest	79150.510	Year	Y
22084	West	691.061	Hour	Y
22388	South	17893.110	Year	Y
23186	Midwest	120195.350	Year	Y
23476	West	95072.750	Year	Y

	case_status
245	Certified
378	Denied
832	Certified
2918	Certified
6439	Denied
6634	Denied
7224	Certified

7281	Denied
7318	Certified
7761	Certified
9872	Certified
11493	Denied
13471	Denied
14022	Certified
14146	Certified
14726	Certified
15600	Denied
15859	Denied
16157	Certified
16883	Denied
17006	Denied
17655	Denied
17844	Denied
17983	Denied
20815	Certified
20984	Certified
21255	Denied
21760	Certified
21944	Certified
22084	Denied
22388	Certified
23186	Certified
23476	Denied

```
In [18]: data['no_of_employees'] = data['no_of_employees'].abs()
updated_negative_rows = data['no_of_employees']
print(updated_negative_rows)
#once the negative vals are found, get the absolute value of them converting them to positive vals
```

0	14513
1	2412
2	44444
3	98
4	1082
	...
25475	2601
25476	3274
25477	1121
25478	1918
25479	3195

Name: no_of_employees, Length: 25480, dtype: int64

Let's check the count of each unique category in each of the categorical variables

```
In [19]: cat_col = data.select_dtypes(include='category').columns
```

```
for col in cat_col:  
    print(f"Value distribution for '{col}':")  
    print(data[col].value_counts(normalize=True))  
    print("-" * 50)
```

#'continent': The majority of the data is from Asia (66.2%), with Europe and North America having smaller proportions.

#'education_of_employee': The dataset is fairly balanced between Bachelor's (40.2%) and Master's (37.8%) degrees, with fewer employees holding a Doctorate (8.6%).

#'has_job_experience': Most individuals have job experience (58.1%), with 41.9% lacking experience.

#'requires_job_training': The majority (88.4%) do not require job training, while only 11.6% do.

#'region_of_employment': Employment is fairly distributed between the Northeast (28.2%), South (27.5%), and West (25.8%), with a small proportion in the Midwest and Island regions.

#'unit_of_wage': Most employees are paid annually (90.1%), with a small fraction paid hourly or weekly.

#'full_time_position': The majority of employees hold full time positions (89.4%).

#'case_status': The majority of cases are certified (66.8%), with a smaller proportion being denied (33.2%).

```

Value distribution for 'case_id':
case_id
EZYV9999    0.000
EZYV01      0.000
EZYV02      0.000
EZYV03      0.000
EZYV04      0.000
...
EZYV10000   0.000
EZYV1000    0.000
EZYV100     0.000
EZYV10      0.000
EZYV09      0.000
Name: proportion, Length: 25480, dtype: float64
-----

```

```

Value distribution for 'continent':
continent
Asia          0.662
Europe        0.146
North America 0.129
South America 0.033
Africa         0.022
Oceania        0.008
Name: proportion, dtype: float64
-----

```

```

Value distribution for 'education_of_employee':
education_of_employee
Bachelor's    0.402
Master's      0.378
High School   0.134
Doctorate     0.086
Name: proportion, dtype: float64
-----

```

```

Value distribution for 'has_job_experience':
has_job_experience
Y    0.581
N    0.419
Name: proportion, dtype: float64
-----

```

```

Value distribution for 'requires_job_training':
requires_job_training
N    0.884
Y    0.116
Name: proportion, dtype: float64
-----

```

```

Value distribution for 'region_of_employment':
region_of_employment
Northeast    0.282
South         0.275
West         0.258
Midwest      0.169
Island        0.015
Name: proportion, dtype: float64
-----

```

```

Value distribution for 'unit_of_wage':
unit_of_wage
Year    0.901

```

```

Hour    0.085
Week    0.011
Month   0.003
Name: proportion, dtype: float64
-----
Value distribution for 'full_time_position':
full_time_position
Y    0.894
N    0.106
Name: proportion, dtype: float64
-----
Value distribution for 'case_status':
case_status
Certified    0.668
Denied       0.332
Name: proportion, dtype: float64
-----

```

In [20]: `data2 = data.copy()`

In [21]: `data2.drop(['case_id'],axis=1,inplace=True)`
#no need for this col, doesnt add value when running models

In [22]: `data2['continent'].unique()`

Out[22]: ['Asia', 'Africa', 'North America', 'Europe', 'South America', 'Oceania']
Categories (6, object): ['Africa', 'Asia', 'Europe', 'North America', 'Oceania', 'South America']

In [23]: `data2['region_of_employment'].unique()`

Out[23]: ['West', 'Northeast', 'South', 'Midwest', 'Island']
Categories (5, object): ['Island', 'Midwest', 'Northeast', 'South', 'West']

In [24]: `data2.head()`

Out[24]:

	continent	education_of_employee	has_job_experience	requires_job_training	no_of_employees
0	Asia	High School	N	N	14513
1	Asia	Master's	Y	N	2412
2	Asia	Bachelor's	N	Y	44444
3	Asia	Bachelor's	N	N	98
4	Africa	Master's	Y	N	1082

In [25]: `data2['education_of_employee'].unique()` *#'prevailing_wage', 'yr_of_estab', 'no_of_employees', 'has_job_experience', 'requires_job_training', 'full_time_position', 'case_status', unit_of_wage*

Out[25]: ['High School', 'Master's', 'Bachelor's', 'Doctorate']
Categories (4, object): ['Bachelor's', 'Doctorate', 'High School', 'Master's']


```
In [26]: data2['unit_of_wage'].unique()
```

```
Out[26]: ['Hour', 'Year', 'Week', 'Month']
Categories (4, object): ['Hour', 'Month', 'Week', 'Year']
```

```
In [27]: data2['education_of_employee'].replace({
    'High School': 1,
    'Bachelor's': 2,
    'Master's': 3,
    'Doctorate': 4
}, inplace=True)

data2['has_job_experience'].replace({'Y': 1, 'N': 0}, inplace=True)
data2['requires_job_training'].replace({'Y': 1, 'N': 0}, inplace=True)
data2['full_time_position'].replace({'Y': 1, 'N': 0}, inplace=True)
data2['case_status'].replace({'Certified': 1, 'Denied': 0}, inplace=True)

region_mapping = {
    'West': 0,
    'Northeast': 1,
    'South': 2,
    'Midwest': 3,
    'Island': 4
}
data2['region_of_employment'] = data2['region_of_employment'].replace(region_mapping)

continent_mapping = {
    'North America': 0,
    'Europe': 1,
    'South America': 2,
    'Africa': 3,
    'Oceania': 4,
    'Asia': 5
}
data2['continent'] = data2['continent'].replace(continent_mapping)

wage_mapping = {
    'Hour': 1,
    'Month': 2,
    'Week': 3,
    'Year': 4
}
data2['unit_of_wage'] = data2['unit_of_wage'].replace(wage_mapping)

#converting all columns with string vals to one hot encoding for better consumption when running the models
```

Univariate Analysis

```
In [28]: def histogram_boxplot(data, feature, figsize=(15, 10), kde=False, bins=None):
        """
        Boxplot and histogram combined

        data: dataframe
        feature: dataframe column
        figsize: size of figure (default (15,10))
        kde: whether to show the density curve (default False)
        bins: number of bins for histogram (default None)
        """
        f2, (ax_box2, ax_hist2) = plt.subplots(
            nrows=2, # Number of rows of the subplot grid= 2
            sharex=True, # x-axis will be shared among all subplots
            gridspec_kw={"height_ratios": (0.25, 0.75)},
            figsize=figsize,
        ) # creating the 2 subplots
        sns.boxplot(
            data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
        ) # boxplot will be created and a triangle will indicate the mean value o
f the column
        ax_box2.set_title(f'Boxplot of {feature}')
        sns.histplot(
            data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins
        ) if bins else sns.histplot(
            data=data, x=feature, kde=kde, ax=ax_hist2
        ) # For histogram
        ax_hist2.axvline(
            data[feature].mean(), color="green", linestyle="--"
        ) # Add mean to the histogram
        ax_hist2.axvline(
            data[feature].median(), color="black", linestyle="-"
        ) # Add median to the histogram
```

In [29]: *# function to create labeled barplots*

```
def labeled_barplot(data, feature, perc=False, n=None):
    """
    Barplot with percentage at the top

    data: dataframe
    feature: dataframe column
    perc: whether to display percentages instead of count (default is False)
    n: displays the top n category levels (default is None, i.e., display all
    levels)
    """

    total = len(data[feature]) # length of the column
    count = data[feature].nunique()
    if n is None:
        plt.figure(figsize=(count + 1, 5))
    else:
        plt.figure(figsize=(n + 1, 5))

    plt.xticks(rotation=90, fontsize=15)
    ax = sns.countplot(
        data=data,
        x=feature,
        palette="Paired",
        order=data[feature].value_counts().index[:n].sort_values(),
    )
    ax.set_title(f'Boxplot of {feature}')

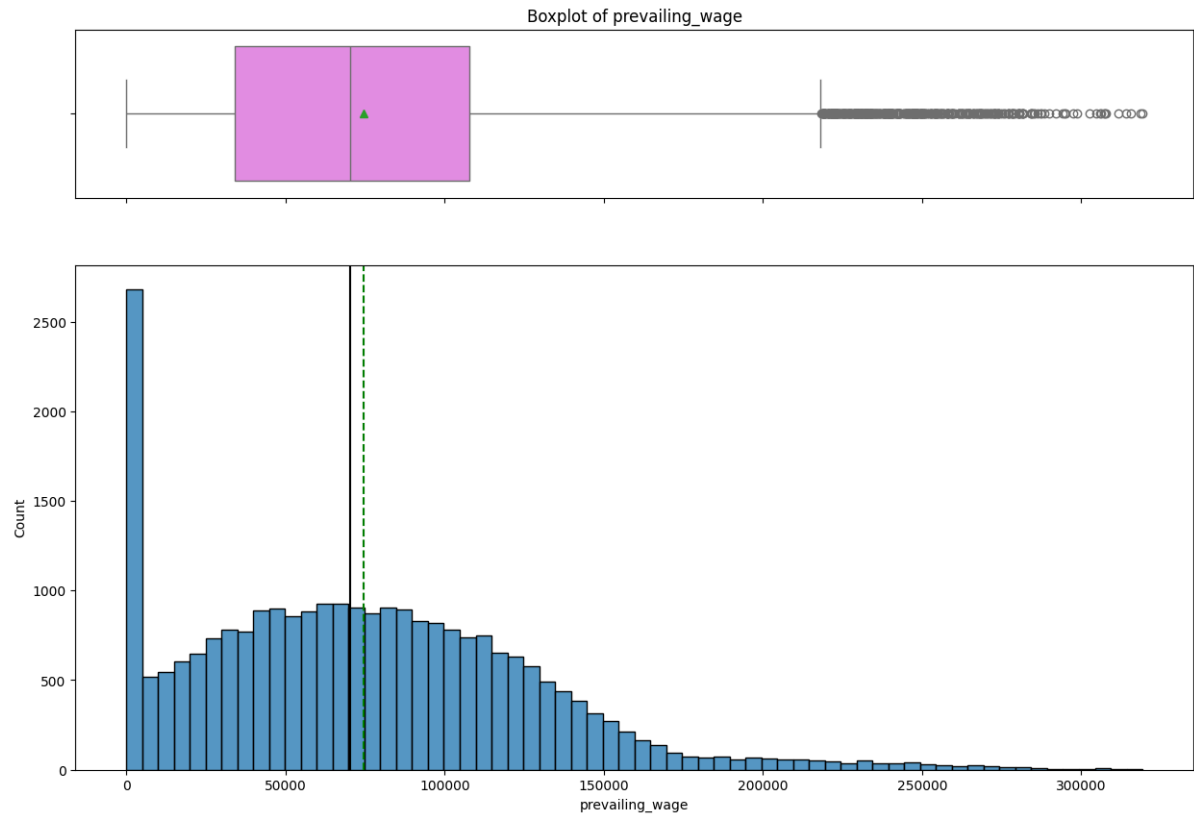
    for p in ax.patches:
        if perc == True:
            label = "{:.1f}%".format(
                100 * p.get_height() / total
            ) # percentage of each class of the category
        else:
            label = p.get_height() # count of each level of the category

        x = p.get_x() + p.get_width() / 2 # width of the plot
        y = p.get_height() # height of the plot

        ax.annotate(
            label,
            (x, y),
            ha="center",
            va="center",
            size=12,
            xytext=(0, 5),
            textcoords="offset points",
        ) # annotate the percentage

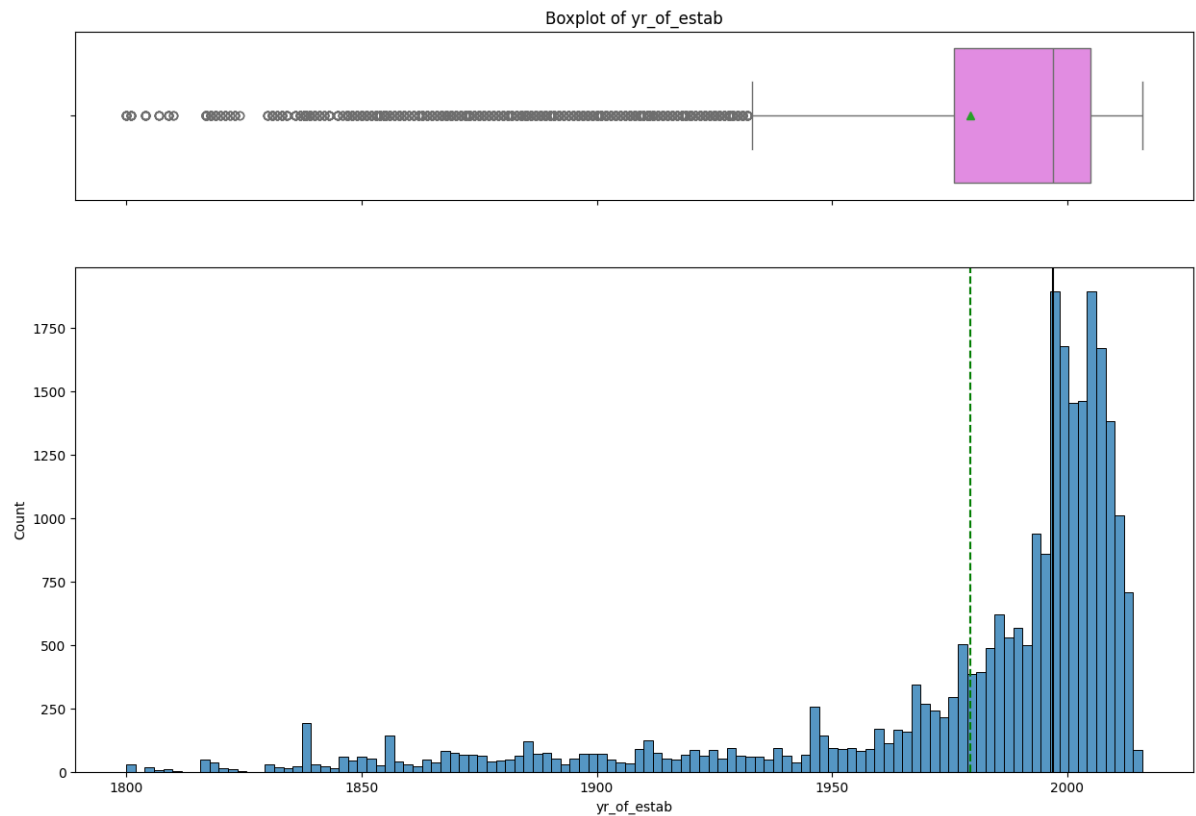
    plt.show() # show the plot
```

```
In [30]: histogram_boxplot(data, 'prevailing_wage')  
#majority of people in the ds dont have a prevailing wage since they are new gr  
ads, but there are a lot of outliers that are making a ton of money  
#avg wage is around 70k or so
```

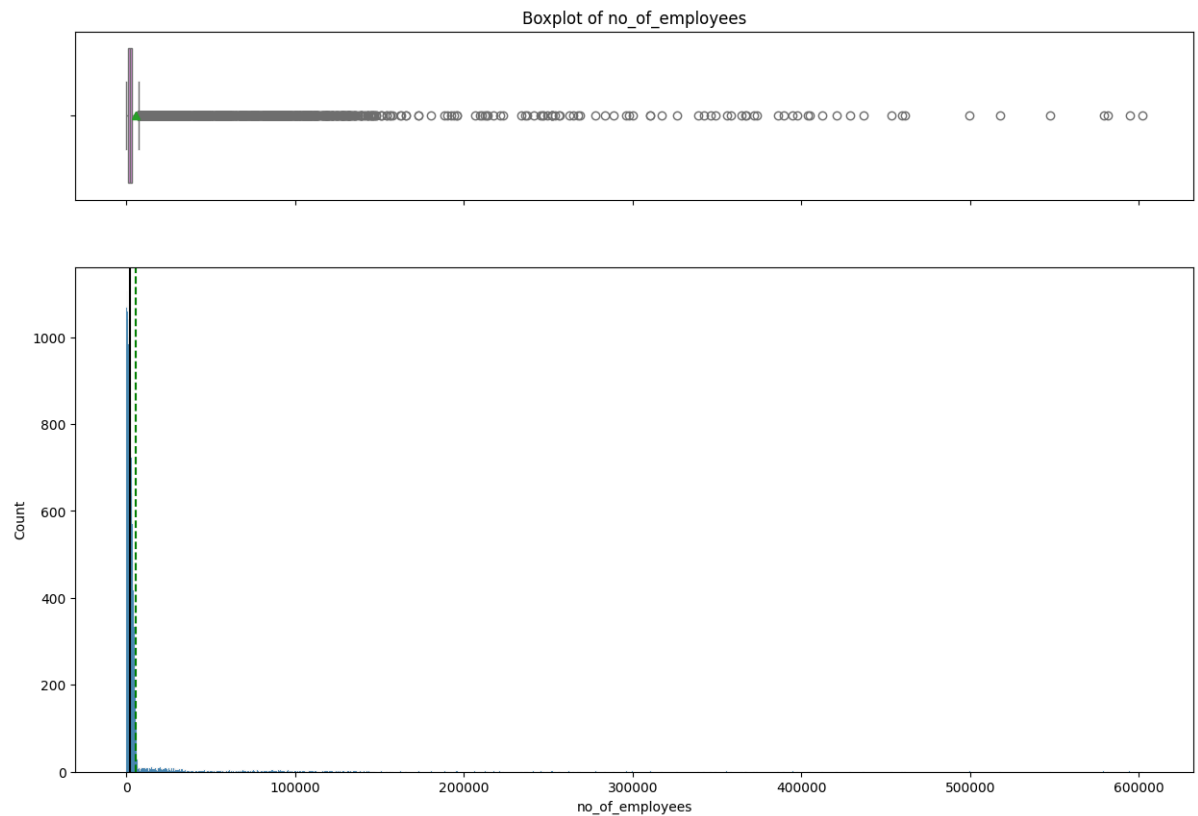


```
In [31]: histogram_boxplot(data, 'yr_of_estab')
```

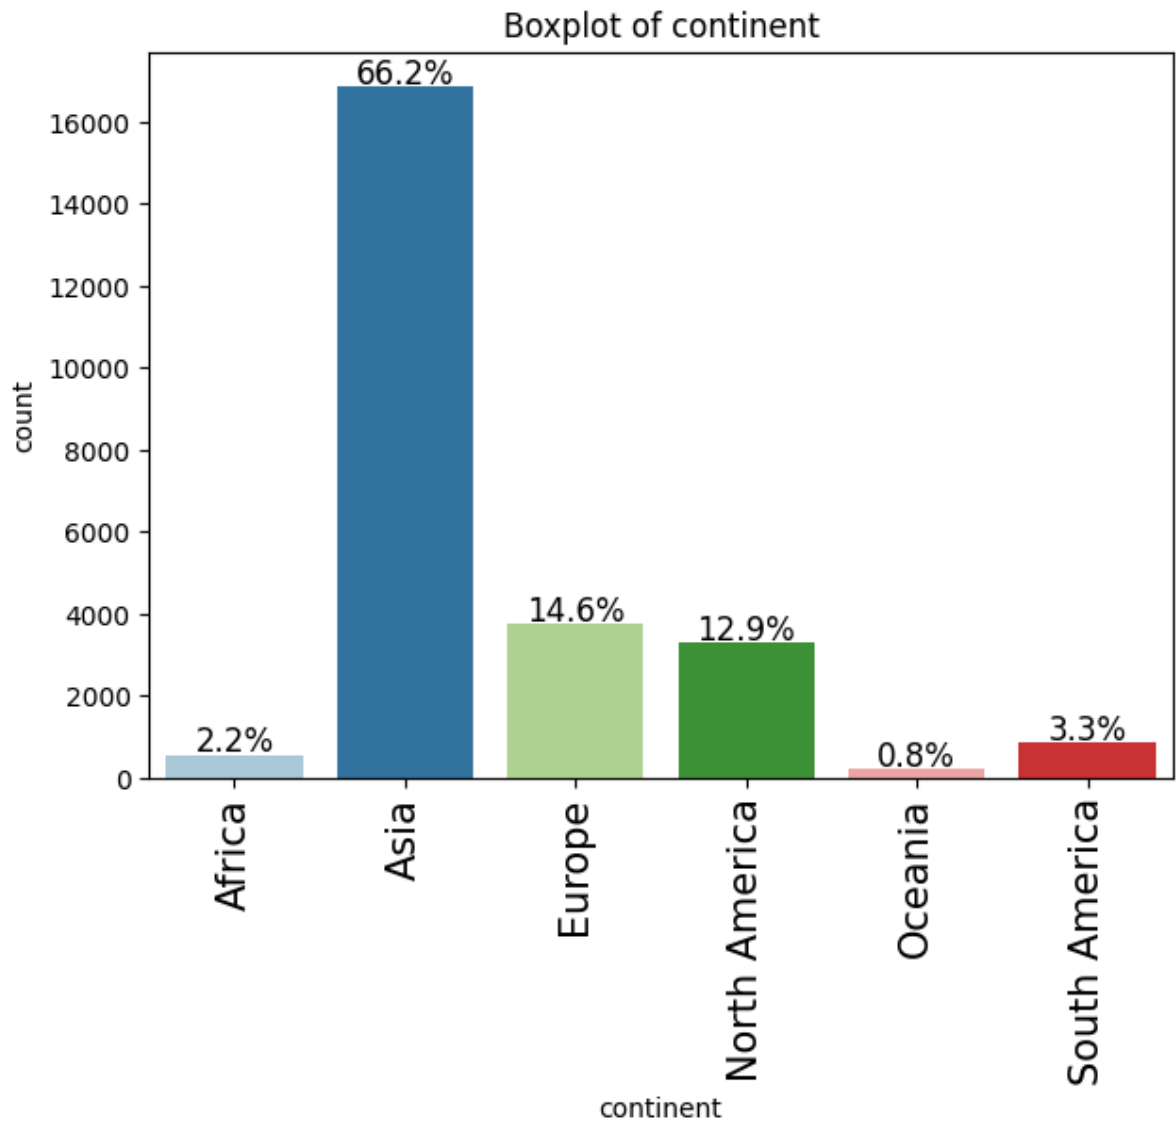
#majority of companies were established around 2000, dot com boom, there are a ton of outliers but since this is yr the math might not be correct



```
In [32]: histogram_boxplot(data, 'no_of_employees')  
#a lot of companies are maybe only a few emps, but others are well established  
so they have more
```

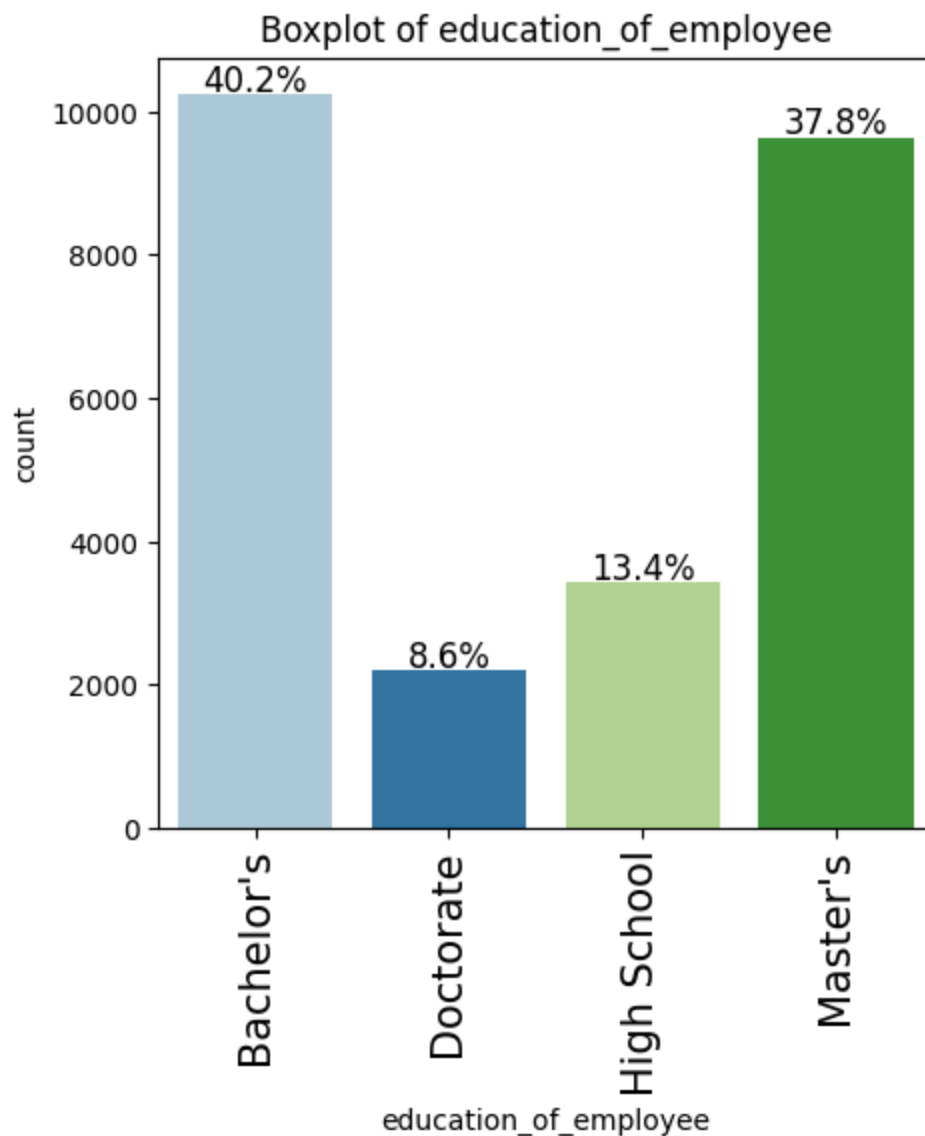


```
In [33]: labeled_barplot(data, 'continent', perc = True)  
#nearly 70% of workers are located in asia
```



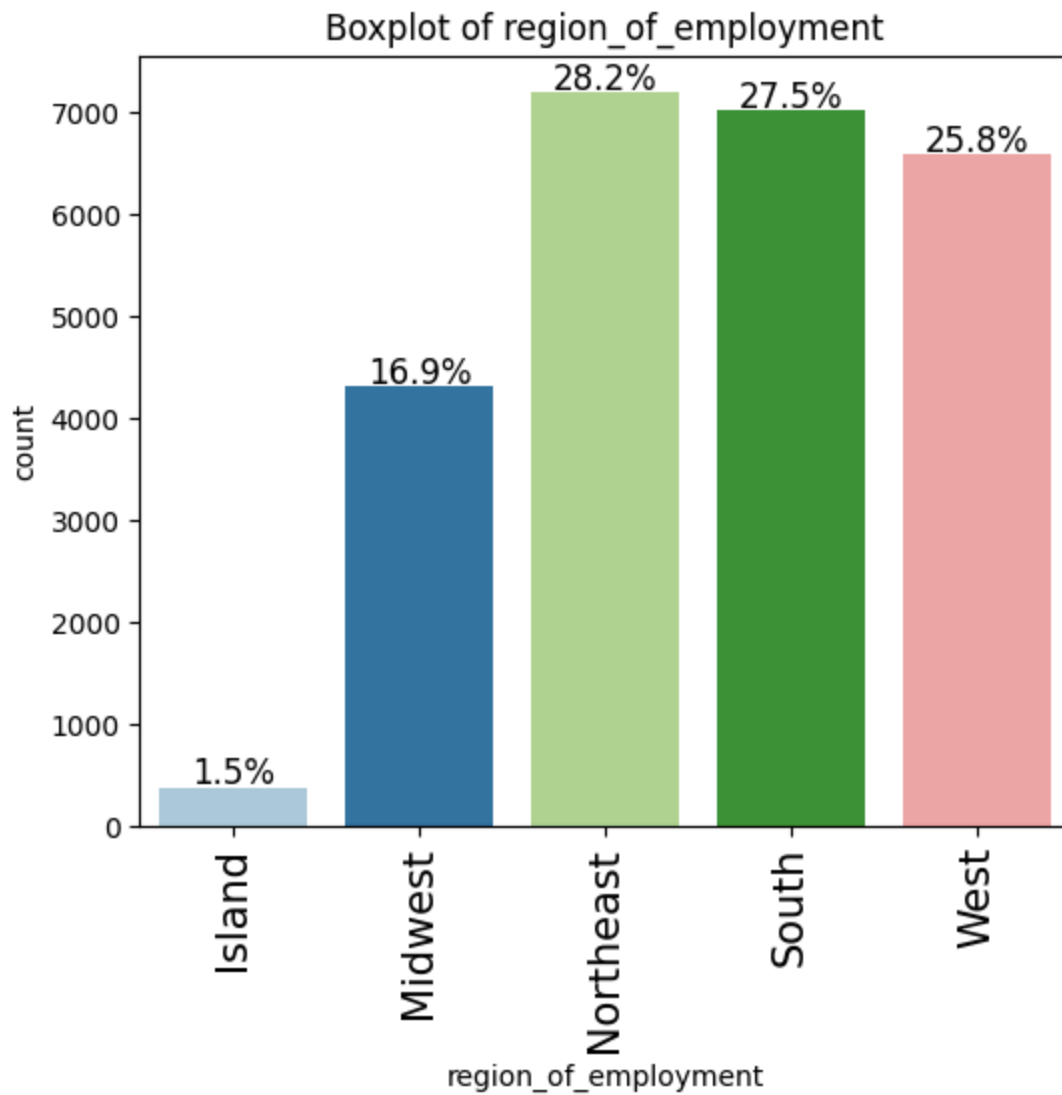
Observations on education of employee

```
In [34]: labeled_barplot(data, 'education_of_employee', perc = True)
#clear split between new grads and BS and veterans and masters holders that are applying for a case
```



Observations on region of employment

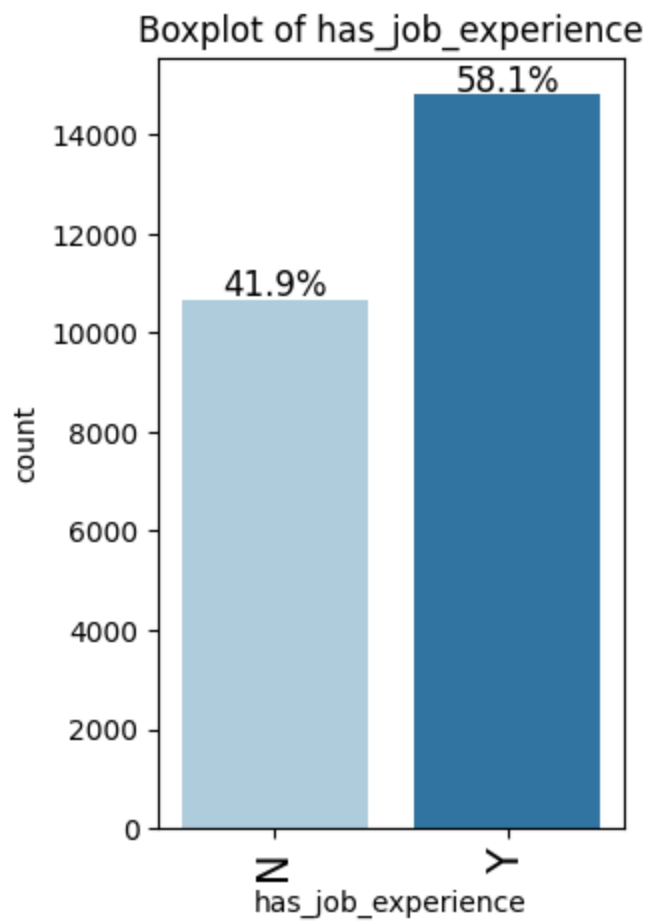
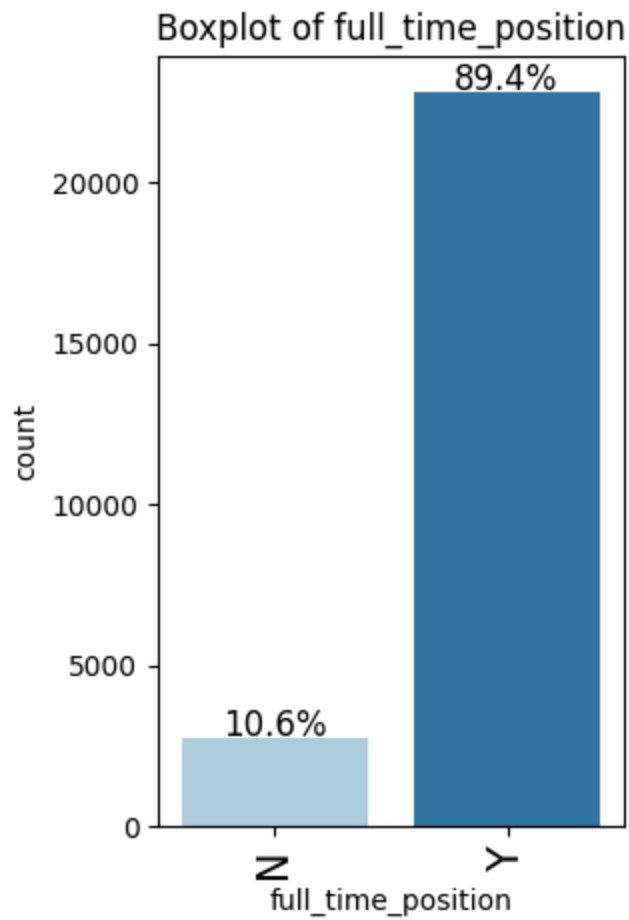

```
In [35]: labeled_barplot(data, 'region_of_employment', perc = True)
#this map is a bit confusing since we dont know for sure which part of the world this is talking about
```

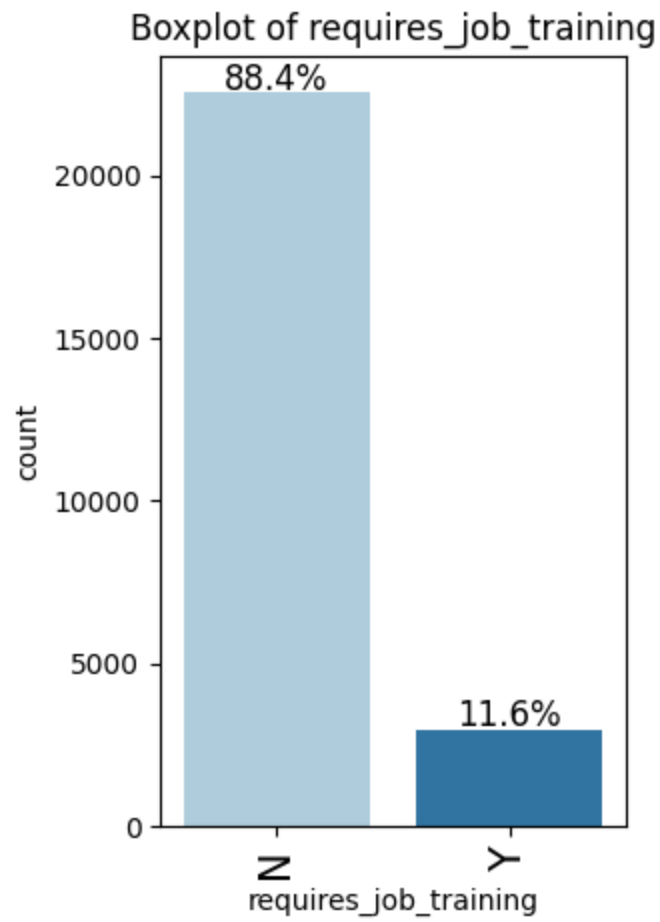


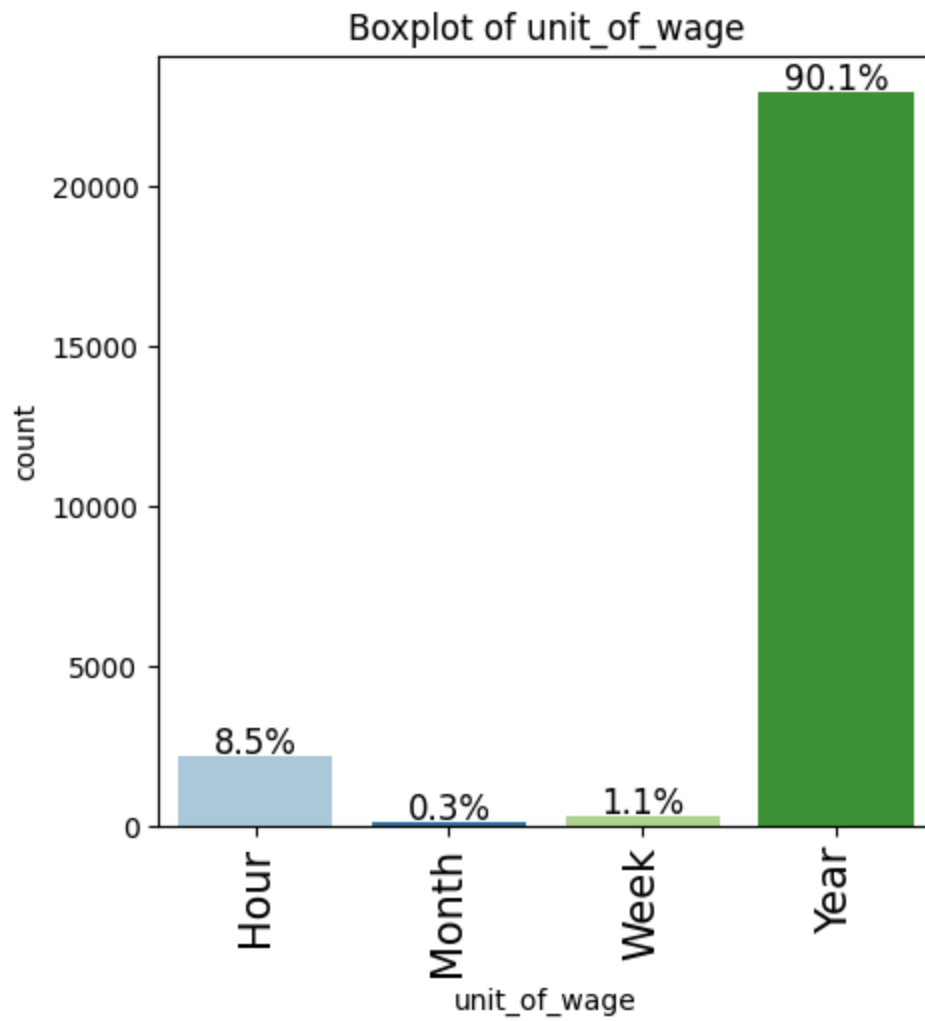
Observations on job experience

```
In [36]: labeled_barplot(data, 'full_time_position', perc = True)
labeled_barplot(data, 'has_job_experience', perc = True)
labeled_barplot(data, 'requires_job_training', perc = True)
labeled_barplot(data, 'unit_of_wage', perc = True)

# 90% of workers have a full time job which is good for getting a case
# a little more than half of the ds has work exp, which means theyll get a cas
e
# 90% of workers dont need training which means that they have prior exp
# 90% of workers have jobs, which means they get a salary, but near 10% are po
ssibly contractors
```

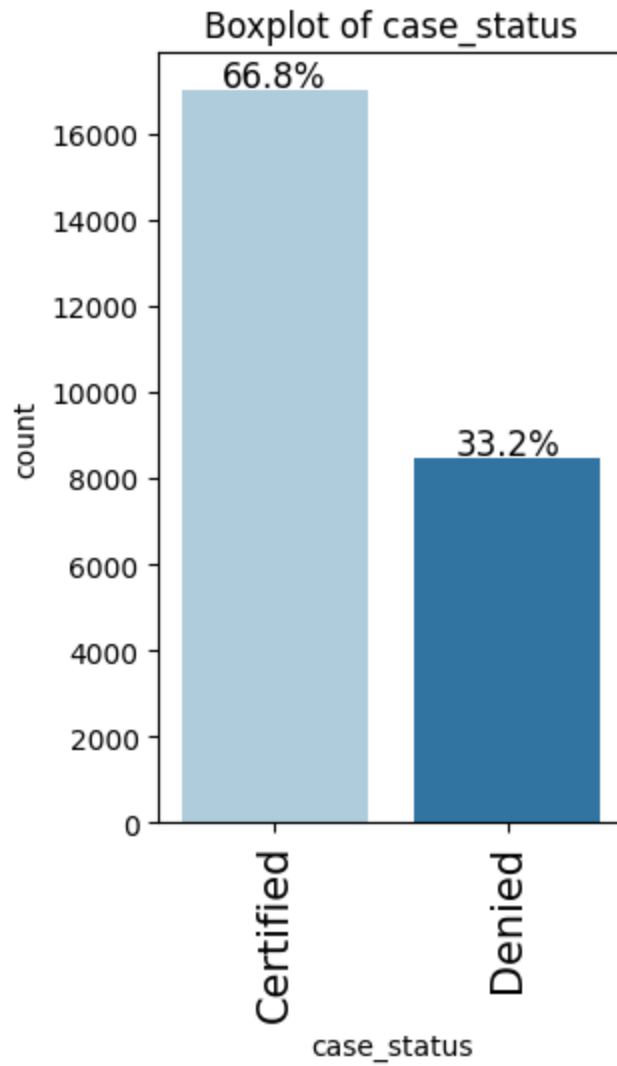






Observations on case status

```
In [37]: labeled_barplot(data, 'case_status', perc = True)  
# 2/3 of workers are getting their case certified
```



Bivariate Analysis

Creating functions that will help us with further analysis.

```

In [38]: ### function to plot distributions wrt target

def distribution_plot_wrt_target(data, predictor, target):

    fig, axs = plt.subplots(2, 2, figsize=(12, 10))

    target_uniq = data[target].unique()

    axs[0, 0].set_title("Distribution of target for target=" + str(target_uniq
[0]))
    sns.histplot(
        data=data[data[target] == target_uniq[0]],
        x=predictor,
        kde=True,
        ax=axs[0, 0],
        color="teal",
        stat="density",
    )

    axs[0, 1].set_title("Distribution of target for target=" + str(target_uniq
[1]))
    sns.histplot(
        data=data[data[target] == target_uniq[1]],
        x=predictor,
        kde=True,
        ax=axs[0, 1],
        color="orange",
        stat="density",
    )

    axs[1, 0].set_title("Boxplot w.r.t target")
    sns.boxplot(data=data, x=target, y=predictor, ax=axs[1, 0], palette="gist_
rainbow")

    axs[1, 1].set_title("Boxplot (without outliers) w.r.t target")
    sns.boxplot(
        data=data,
        x=target,
        y=predictor,
        ax=axs[1, 1],
        showfliers=False,
        palette="gist_rainbow",
    )

    plt.tight_layout()
    plt.show()

```

```

In [39]: def stacked_barplot(data, predictor, target):
        """
        Print the category counts and plot a stacked bar chart

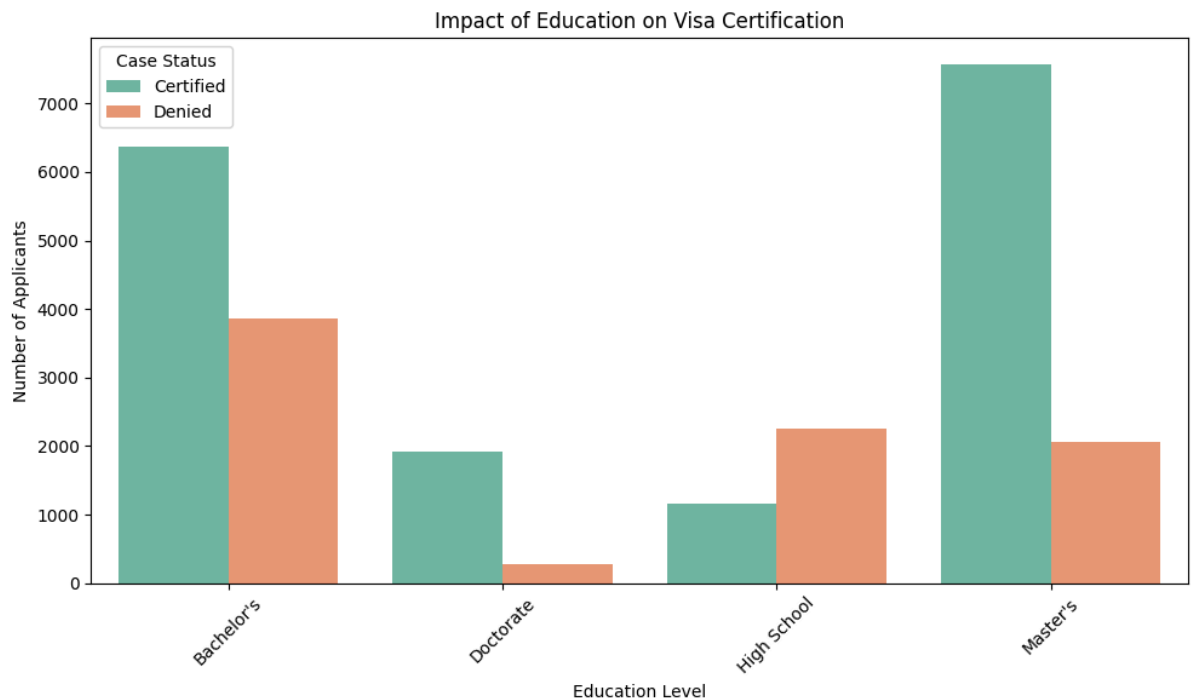
        data: dataframe
        predictor: independent variable
        target: target variable
        """
        count = data[predictor].nunique()
        sorter = data[target].value_counts().index[-1]
        tab1 = pd.crosstab(data[predictor], data[target], margins=True).sort_value
s(
        by=sorter, ascending=False
    )
    print(tab1)
    print("-" * 120)
    tab = pd.crosstab(data[predictor], data[target], normalize="index").sort_v
alues(
        by=sorter, ascending=False
    )
    tab.plot(kind="bar", stacked=True, figsize=(count + 5, 5))
    plt.legend(
        loc="lower left", frameon=False,
    )
    plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
    plt.show()

```

Those with higher education may want to travel abroad for a well-paid job. Let's find out if education has any impact on visa certification


```
In [40]: plt.figure(figsize=(10, 6))
sns.countplot(data=data, x='education_of_employee', hue='case_status', palette='Set2')
plt.title('Impact of Education on Visa Certification')
plt.xlabel('Education Level')
plt.ylabel('Number of Applicants')
plt.legend(title='Case Status')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

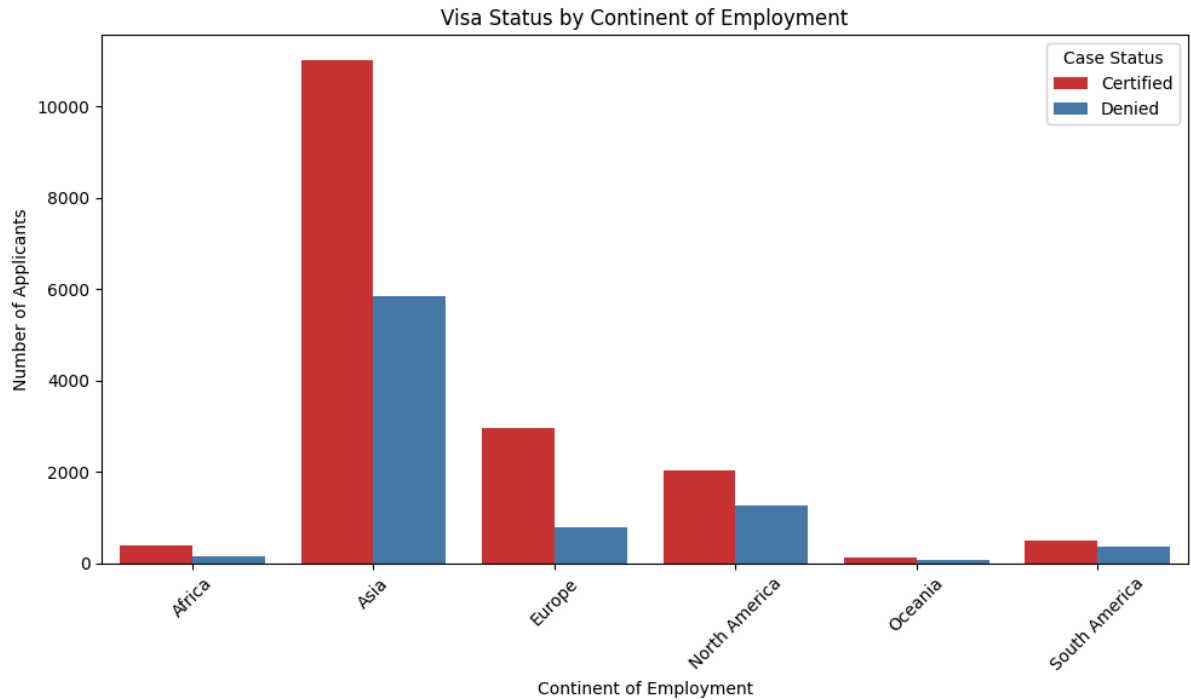
# a majority of workers that have either a beginner education or a seasoned education are getting their cases certified, HS workers probably won't get it
# good portion of phd folks are getting certified
```



Lets' similarly check for the continents and find out how the visa status vary across different continents.

```
In [41]: plt.figure(figsize=(10, 6))
sns.countplot(data=data, x='continent', hue='case_status', palette='Set1')
plt.title('Visa Status by Continent of Employment')
plt.xlabel('Continent of Employment')
plt.ylabel('Number of Applicants')
plt.legend(title='Case Status')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

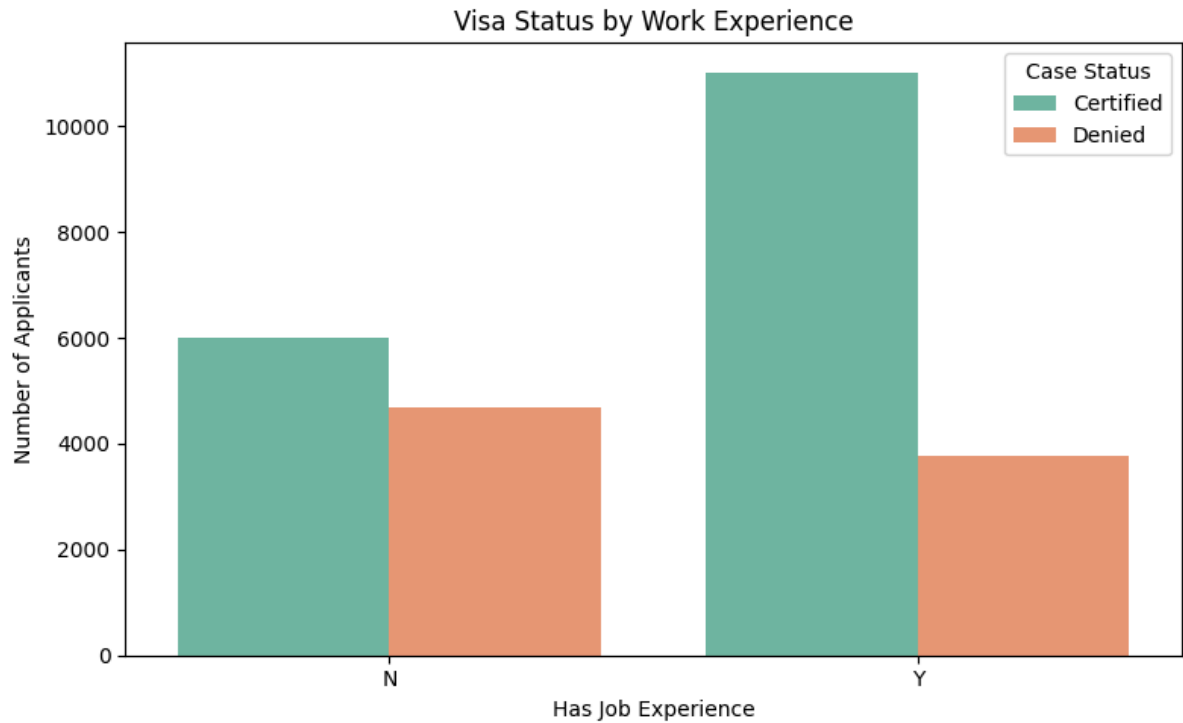
# mainly only folks in asia are getting certified for cases, this seems that t
hey are qualified and better than the rest of the continents
```



Experienced professionals might look abroad for opportunities to improve their lifestyles and career development. Let's see if having work experience has any influence over visa certification

```
In [42]: plt.figure(figsize=(8, 5))
sns.countplot(data=data, x='has_job_experience', hue='case_status', palette='Set2')
plt.title('Visa Status by Work Experience')
plt.xlabel('Has Job Experience')
plt.ylabel('Number of Applicants')
plt.legend(title='Case Status')
plt.tight_layout()
plt.show()

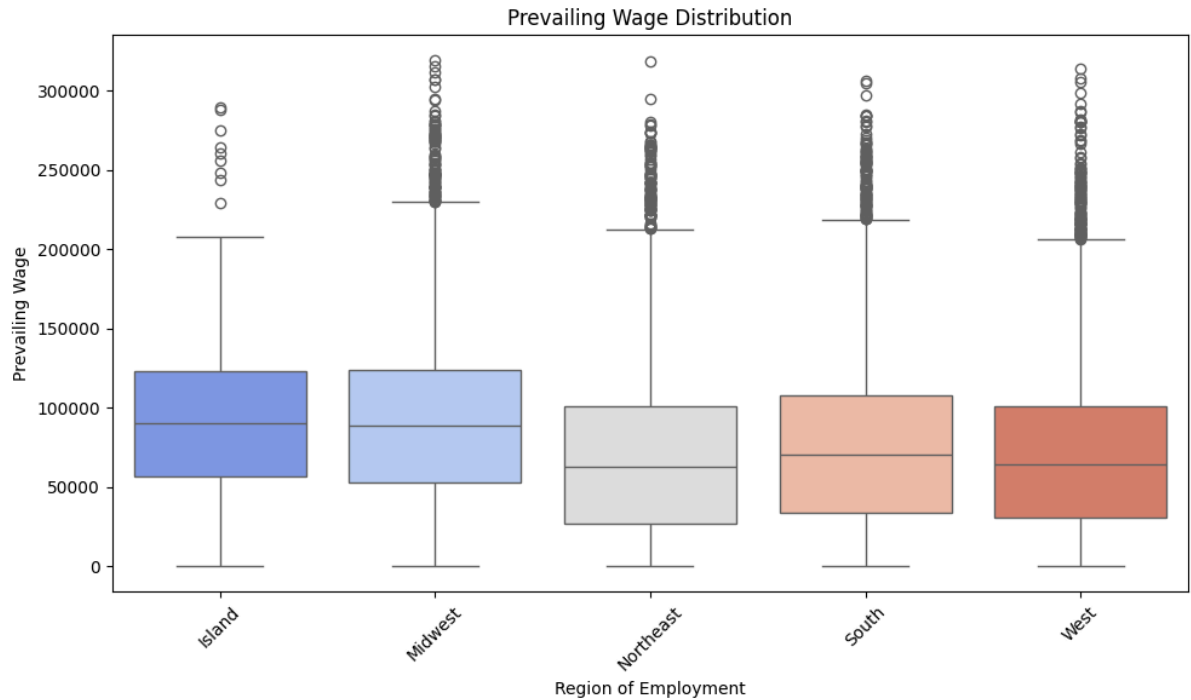
# if workers have job exp then they are more likely to get a visa, which makes sense, we need workers
```



Checking if the prevailing wage is similar across all the regions of the US

```
In [43]: plt.figure(figsize=(10, 6))
sns.boxplot(data=data, x='region_of_employment', y='prevailing_wage', palette
='coolwarm')
plt.title('Prevailing Wage Distribution')
plt.xlabel('Region of Employment')
plt.ylabel('Prevailing Wage')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

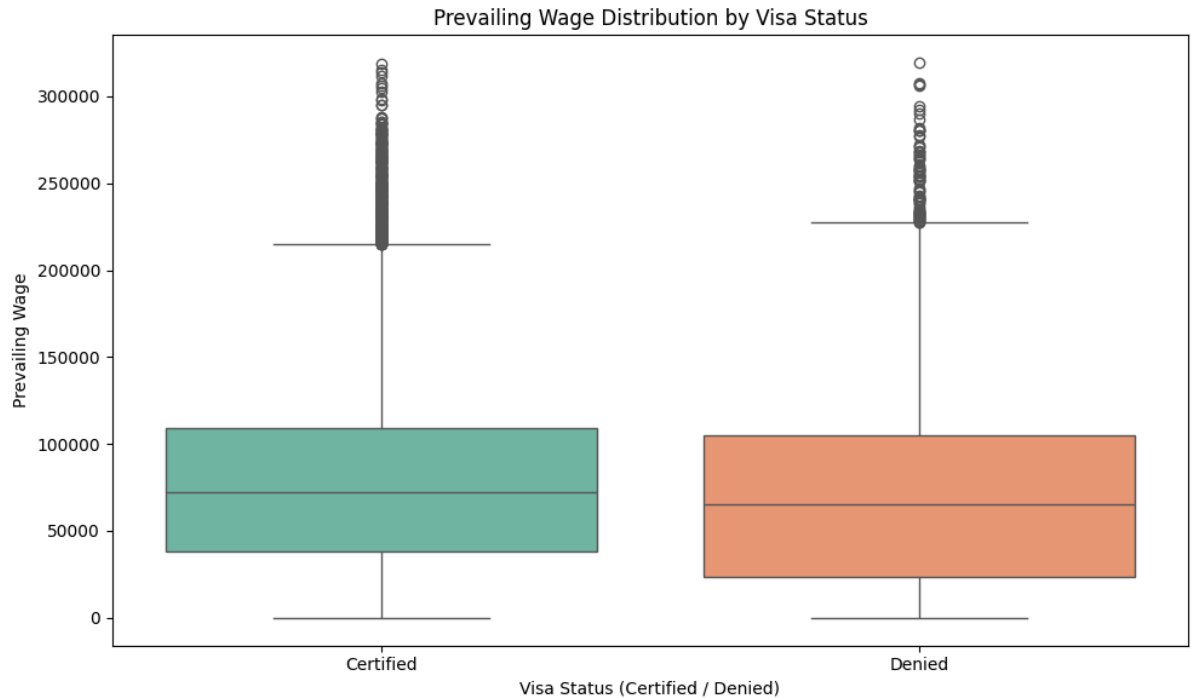
# the midwest region has a lot of outliers in the higher end of wages, but nam
ing scheme is a bit tough to understand which part of world exactly
```



The US government has established a prevailing wage to protect local talent and foreign workers. Let's analyze the data and see if the visa status changes with the prevailing wage

```
In [44]: plt.figure(figsize=(10, 6))
sns.boxplot(data=data, x='case_status', y='prevailing_wage', palette='Set2')
plt.title('Prevailing Wage Distribution by Visa Status')
plt.xlabel('Visa Status (Certified / Denied)')
plt.ylabel('Prevailing Wage')
plt.tight_layout()
plt.show()
```

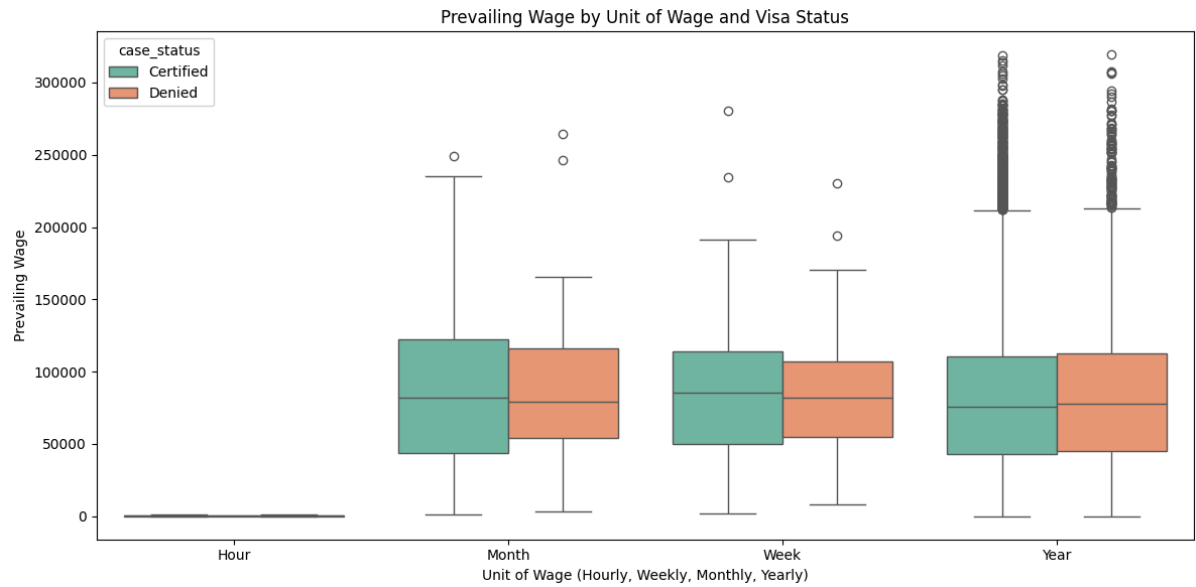
there is a higher baseline for wages for people that have visas and they have more outliers, 50% of them make slightly more than the denied folks



The prevailing wage has different units (Hourly, Weekly, etc). Let's find out if it has any impact on visa applications getting certified.

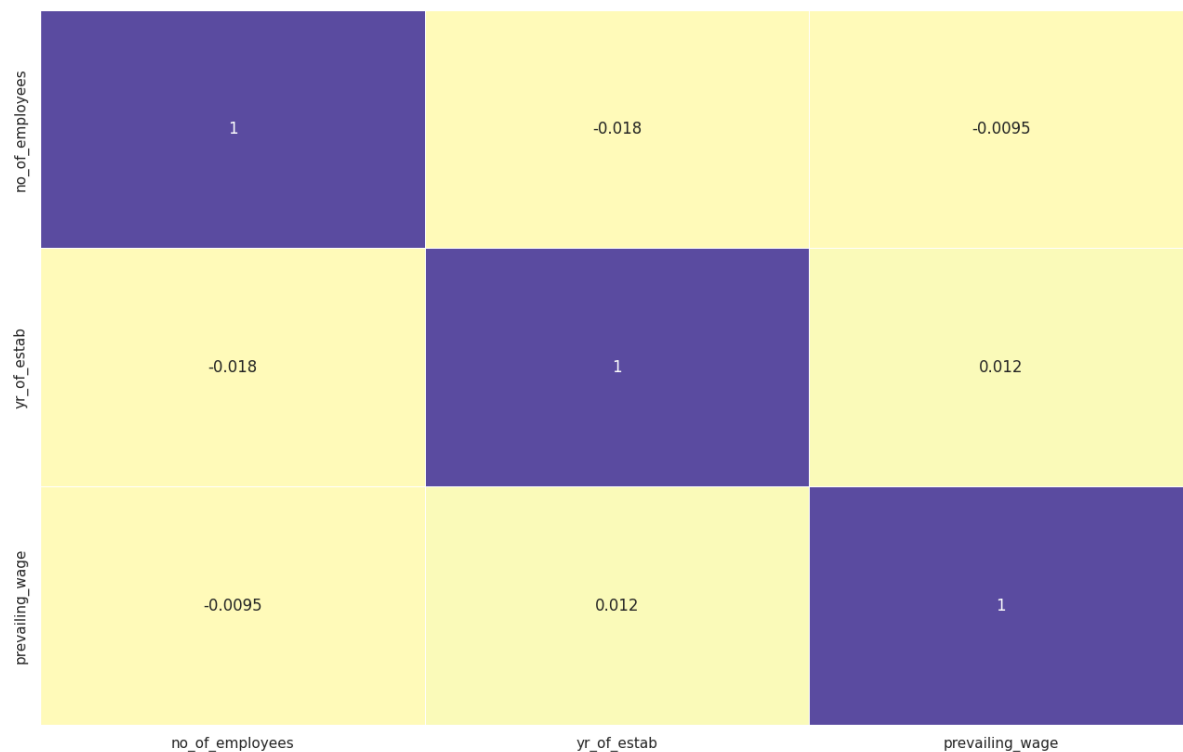
```
In [45]: plt.figure(figsize=(12, 6))
sns.boxplot(data=data, x='unit_of_wage', y='prevailing_wage', hue='case_status', palette='Set2')
plt.title('Prevailing Wage by Unit of Wage and Visa Status')
plt.xlabel('Unit of Wage (Hourly, Weekly, Monthly, Yearly)')
plt.ylabel('Prevailing Wage')
plt.tight_layout()
plt.show()

# hourly wages dont exist basically, yearly wages is split essentially since half are certified and other half is denied
```



```
In [46]: sns.set(rc={'figure.figsize': (16, 10)})

sns.heatmap(
    data.corr(numeric_only=True),
    annot=True,
    linewidths=0.5,
    center=0,
    cbar=False,
    cmap="Spectral"
)
plt.show()
```



Data Pre-processing

Outlier Check

```
In [47]: data_encoded = data2.copy()

for col in data_encoded.select_dtypes(include='category').columns:
    data_encoded[col] = data_encoded[col].cat.codes

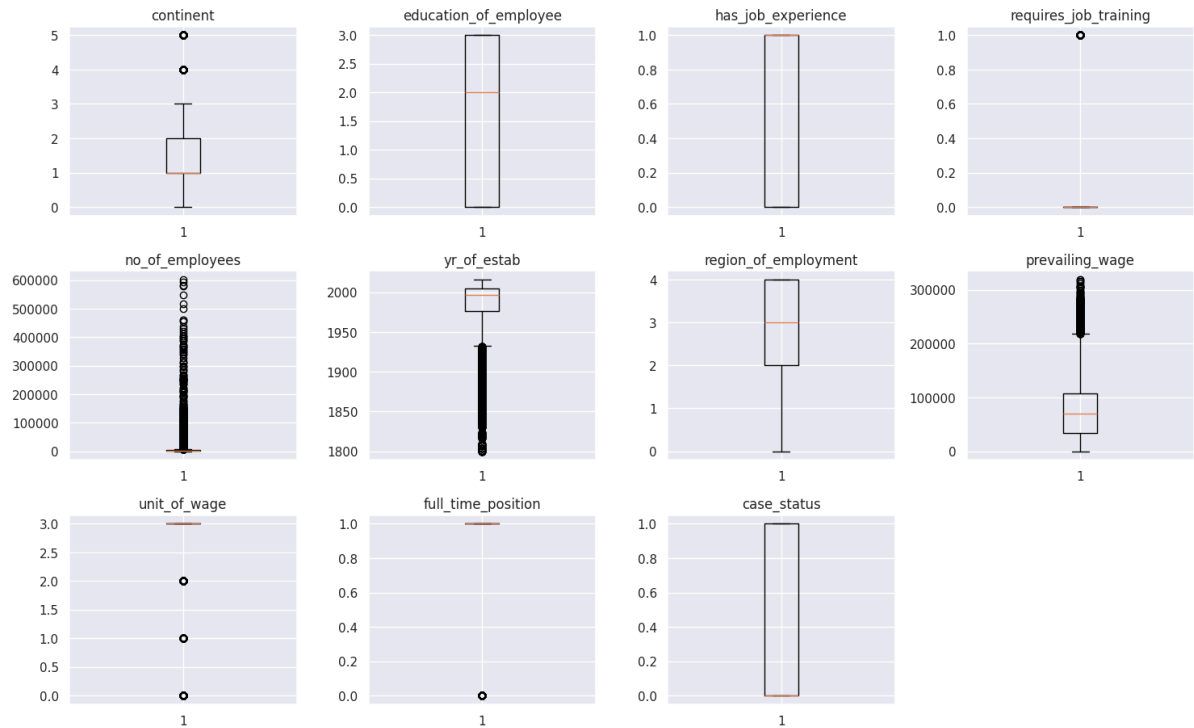
numeric_columns = data_encoded.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(15, 12))

for i, variable in enumerate(numeric_columns):
    plt.subplot(4, 4, i + 1)
```

```
plt.boxplot(data_encoded[variable], whis=1.5)
plt.tight_layout()
plt.title(variable)
```

```
plt.show()
```



Data Preparation for modeling

```
In [48]: X = data2.drop(["case_status"], axis=1)
         y = data2["case_status"]
```

```
# dropping the target variable
```

```
In [49]: X_temp, X_test, y_temp, y_test = train_test_split(
         X, y, test_size=0.3, random_state=1, stratify=y
         )

# then we split the temporary set into train and validation

X_train, X_val, y_train, y_val = train_test_split(
    X_temp, y_temp, test_size=0.25, random_state=1, stratify=y_temp
)
print(X_train.shape, X_val.shape, X_test.shape)
```

```
# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=1, shuffle=True)
```

```
(13377, 10) (4459, 10) (7644, 10)
```



```
In [50]: X_train.shape, X_val.shape
```

```
Out[50]: ((13377, 10), (4459, 10))
```

```
In [51]: import pandas as pd
from sklearn.impute import SimpleImputer

# Get list of categorical and numerical columns
cat_cols = list(X_train.select_dtypes(include='category').columns)
num_cols = list(X_train.select_dtypes(include=['int', 'float']).columns)

# Impute categorical columns
cat_imputer = SimpleImputer(strategy='most_frequent')
X_train[cat_cols] = cat_imputer.fit_transform(X_train[cat_cols])
X_val[cat_cols] = cat_imputer.transform(X_val[cat_cols])
X_test[cat_cols] = cat_imputer.transform(X_test[cat_cols])

# Impute numerical columns
num_imputer = SimpleImputer(strategy='mean')
X_train[num_cols] = num_imputer.fit_transform(X_train[num_cols])
X_val[num_cols] = num_imputer.transform(X_val[num_cols])
X_test[num_cols] = num_imputer.transform(X_test[num_cols])
data2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25480 entries, 0 to 25479
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   continent                            25480 non-null  category
1   education_of_employee                25480 non-null  category
2   has_job_experience                    25480 non-null  category
3   requires_job_training                 25480 non-null  category
4   no_of_employees                      25480 non-null  int64
5   yr_of_estab                          25480 non-null  int64
6   region_of_employment                 25480 non-null  category
7   prevailing_wage                      25480 non-null  float64
8   unit_of_wage                         25480 non-null  category
9   full_time_position                   25480 non-null  category
10  case_status                           25480 non-null  category
dtypes: category(8), float64(1), int64(2)
memory usage: 797.7 KB
```

```
In [52]: # Checking that no column has missing values in train, validation or test sets
print(X_train.isna().sum())
print("-" * 30)
print(X_val.isna().sum())
print("-" * 30)
print(X_test.isna().sum())
```

```
continent          0
education_of_employee  0
has_job_experience  0
requires_job_training  0
no_of_employees     0
yr_of_estab         0
region_of_employment  0
prevailing_wage     0
unit_of_wage        0
full_time_position  0
dtype: int64
```

```
-----
continent          0
education_of_employee  0
has_job_experience  0
requires_job_training  0
no_of_employees     0
yr_of_estab         0
region_of_employment  0
prevailing_wage     0
unit_of_wage        0
full_time_position  0
dtype: int64
```

```
-----
continent          0
education_of_employee  0
has_job_experience  0
requires_job_training  0
no_of_employees     0
yr_of_estab         0
region_of_employment  0
prevailing_wage     0
unit_of_wage        0
full_time_position  0
dtype: int64
```

```
In [53]: data2.isnull().sum()
```

```
Out[53]:
```

```

continent 0
education_of_employee 0
has_job_experience 0
requires_job_training 0
no_of_employees 0
yr_of_estab 0
region_of_employment 0
prevailing_wage 0
unit_of_wage 0
full_time_position 0
case_status 0

```

```
dtype: int64
```

Model Building

Model Evaluation Criterion

```
In [54]: # defining a function to compute different metrics to check performance of a c
lassification model built using sklearn

def model_performance_classification_sklearn(model, predictors, target):
    """
    Function to compute different metrics to check classification model perfor
    mance

    model: classifier
    predictors: independent variables
    target: dependent variable
    """

    # predicting using the independent variables
    pred = model.predict(predictors)

    acc = accuracy_score(target, pred) # to compute Accuracy
    recall = recall_score(target, pred) # to compute Recall
    precision = precision_score(target, pred) # to compute Precision
    f1 = f1_score(target, pred) # to compute F1-score

    # creating a dataframe of metrics
```

```

df_perf = pd.DataFrame(
    {"Accuracy": acc, "Recall": recall, "Precision": precision, "F1": f
1,},
    index=[0],
)

return df_perf

```

```

In [55]: def confusion_matrix_sklearn(model, predictors, target):
        """
        To plot the confusion_matrix with percentages

        model: classifier
        predictors: independent variables
        target: dependent variable
        """

        y_pred = model.predict(predictors)
        cm = confusion_matrix(target, y_pred)
        labels = np.asarray(
            [
                ["{0:0.0f}".format(item) + "\n{0:.2%}".format(item / cm.flatten().
sum())]
                for item in cm.flatten()
            ]
        ).reshape(2, 2)

        plt.figure(figsize=(6, 4))
        sns.heatmap(cm, annot=labels, fmt="")
        plt.ylabel("True label")
        plt.xlabel("Predicted label")

```

Defining scorer to be used for cross-validation and hyperparameter tuning

We are now done with pre-processing and evaluation criterion, so let's start building the model.

Model building with original data

```
In [56]: models = []

models.append(("Bagging", BaggingClassifier(estimator=DecisionTreeClassifier(r
random_state=1, class_weight='balanced'), random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1, class_w
eight='balanced')))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("dtree", DecisionTreeClassifier(random_state=1, class_weight='b
alanced')))

print("\nTraining Performance:\n")
for name, model in models:
    model.fit(X_train, y_train)
    scores = recall_score(y_train, model.predict(X_train))
    print("{}: {}".format(name, scores))

print("\nValidation Performance:\n")
for name, model in models:
    model.fit(X_train, y_train)
    scores_val = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores_val))

# adding the 5 models and running the training and validation tests on it, thi
s is right out of the gate before anything is done to it
```

Training Performance:

Bagging: 0.9863458310016788
Random forest: 1.0
GBM: 0.8846110800223839
Adaboost: 0.8878567431449357
dtree: 1.0

Validation Performance:

Bagging: 0.7991940899932841
Random forest: 0.8707186030893217
GBM: 0.8935527199462726
Adaboost: 0.8989254533243788
dtree: 0.7635997313633311

```
In [57]: print("\nTraining and Validation Performance Difference:\n")

for name, model in models:
    model.fit(X_train, y_train)
    scores_train = recall_score(y_train, model.predict(X_train))
    scores_val = recall_score(y_val, model.predict(X_val))
    difference1 = scores_train - scores_val
    print("{}: Training Score: {:.4f}, Validation Score: {:.4f}, Difference: {:.4f}".format(name, scores_train, scores_val, difference1))

# getting the diff between the 2 tests to see how they stack up against each other and also to see room for improvement.
```

Training and Validation Performance Difference:

Bagging: Training Score: 0.9863, Validation Score: 0.7992, Difference: 0.1872
Random forest: Training Score: 1.0000, Validation Score: 0.8707, Difference: 0.1293
GBM: Training Score: 0.8846, Validation Score: 0.8936, Difference: -0.0089
Adaboost: Training Score: 0.8879, Validation Score: 0.8989, Difference: -0.0111
dtree: Training Score: 1.0000, Validation Score: 0.7636, Difference: 0.2364

Model Building with oversampled data

```
In [58]: print("Before Oversampling, counts of label 'Yes': {}".format(sum(y_train ==
1)))
print("Before Oversampling, counts of label 'No': {} \n".format(sum(y_train ==
0)))

sm = SMOTE(
    sampling_strategy=1, k_neighbors=5, random_state=1
) # Synthetic Minority Over Sampling Technique
X_train_over, y_train_over = sm.fit_resample(X_train, y_train)

print("After Oversampling, counts of label 'Yes': {}".format(sum(y_train_over
== 1)))
print("After Oversampling, counts of label 'No': {} \n".format(sum(y_train_ove
r == 0)))

print("After Oversampling, the shape of train_X: {}".format(X_train_over.shap
e))
print("After Oversampling, the shape of train_y: {} \n".format(y_train_over.sh
ape))

# running SMOTE on the dataset for over/under sampling to see if the ds change
s at all, only the over changes for NO after running it which means some folks
were on the border
```

Before Oversampling, counts of label 'Yes': 8935
Before Oversampling, counts of label 'No': 4442

After Oversampling, counts of label 'Yes': 8935
After Oversampling, counts of label 'No': 8935

After Oversampling, the shape of train_X: (17870, 10)
After Oversampling, the shape of train_y: (17870,)

```
In [59]: X_train_over.isnull().sum()
```

```
Out[59]:
```

	0
continent	0
education_of_employee	0
has_job_experience	0
requires_job_training	0
no_of_employees	0
yr_of_estab	0
region_of_employment	0
prevailing_wage	0
unit_of_wage	0
full_time_position	0

```
dtype: int64
```



```
In [60]: models = [] # Empty list to store all the models

# Appending models into the list
models.append(("Bagging", BaggingClassifier(estimator=DecisionTreeClassifier(r
random_state=1, class_weight='balanced'), random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1, class_w
eight='balanced'))))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("dtree", DecisionTreeClassifier(random_state=1, class_weight='b
alanced'))))

print("\n" "Training Performance:" "\n")
for name, model in models:
    model.fit(X_train_over, y_train_over)
    scores = recall_score(y_train_over, model.predict(X_train_over))
    print("{}: {}".format(name, scores))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train_over, y_train_over)
    scores = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores))

# running the validation tests on the OG data to see how the data reacts
# GBM and Ada models did well and increased their numbers which means they gen
eralized well
```

Training Performance:

Bagging: 0.9748181309457191
Random forest: 1.0
GBM: 0.7898153329602686
Adaboost: 0.7838836038052602
dtree: 1.0

Validation Performance:

Bagging: 0.7374076561450638
Random forest: 0.7971793149764943
GBM: 0.7978509066487576
Adaboost: 0.7931497649429147
dtree: 0.708865010073875

```
In [61]: print("\nTraining and Validation Performance Difference:\n")

for name, model in models:
    model.fit(X_train_over, y_train_over)
    scores_train = recall_score(y_train_over, model.predict(X_train_over))
    scores_val = recall_score(y_val, model.predict(X_val))
    difference2 = scores_train - scores_val
    print("{}: Training Score: {:.4f}, Validation Score: {:.4f}, Difference: {:.4f}".format(name, scores_train, scores_val, difference2))
```

Training and Validation Performance Difference:

Bagging: Training Score: 0.9748, Validation Score: 0.7374, Difference: 0.2374

Random forest: Training Score: 1.0000, Validation Score: 0.7972, Difference: 0.2028

GBM: Training Score: 0.7898, Validation Score: 0.7979, Difference: -0.0080

Adaboost: Training Score: 0.7839, Validation Score: 0.7931, Difference: -0.0093

dtree: Training Score: 1.0000, Validation Score: 0.7089, Difference: 0.2911

Model Building with undersampled data

```
In [62]: rus = RandomUnderSampler(random_state=1)
X_train_un, y_train_un = rus.fit_resample(X_train, y_train)
```

```
In [63]: print("Before Under Sampling, counts of label 'Yes': {}".format(sum(y_train == 1)))
print("Before Under Sampling, counts of label 'No': {} \n".format(sum(y_train == 0)))

print("After Under Sampling, counts of label 'Yes': {}".format(sum(y_train_un == 1)))
print("After Under Sampling, counts of label 'No': {} \n".format(sum(y_train_un == 0)))

print("After Under Sampling, the shape of train_X: {}".format(X_train_un.shape))
print("After Under Sampling, the shape of train_y: {} \n".format(y_train_un.shape))
```

Before Under Sampling, counts of label 'Yes': 8935

Before Under Sampling, counts of label 'No': 4442

After Under Sampling, counts of label 'Yes': 4442

After Under Sampling, counts of label 'No': 4442

After Under Sampling, the shape of train_X: (8884, 10)

After Under Sampling, the shape of train_y: (8884,)

```
In [64]: models = [] # Empty list to store all the models

# Appending models into the list
models.append(("Bagging", BaggingClassifier(estimator=DecisionTreeClassifier(r
random_state=1, class_weight='balanced'), random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1, class_w
eight='balanced'))))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("dtree", DecisionTreeClassifier(random_state=1, class_weight='b
alanced'))))

print("\n" "Training Performance:" "\n")
for name, model in models:
    model.fit(X_train_un, y_train_un)
    scores = recall_score(y_train_un, model.predict(X_train_un))
    print("{}: {}".format(name, scores))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train_un, y_train_un)
    scores = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores))

# The high training performance scores could indicate overfitting, where the m
odel is too closely aligned with the training data and might not generalize we
ll
# Models like GBM and Random Forest show relatively good validation performanc
e, which suggests they generalize better than others
```

Training Performance:

Bagging: 0.9669067987393066
Random forest: 1.0
GBM: 0.7478613237280505
Adaboost: 0.7014858171994597
dtree: 1.0

Validation Performance:

Bagging: 0.6229012760241773
Random forest: 0.7081934184016119
GBM: 0.7437877770315648
Adaboost: 0.6984553391537945
dtree: 0.6437206178643384

```
In [65]: print("\nTraining and Validation Performance Difference:\n")

for name, model in models:
    model.fit(X_train_un, y_train_un)
    scores_train = recall_score(y_train_un, model.predict(X_train_un))
    scores_val = recall_score(y_val, model.predict(X_val))
    difference3 = scores_train - scores_val
    print("{}: Training Score: {:.4f}, Validation Score: {:.4f}, Difference: {:.4f}".format(name, scores_train, scores_val, difference3))
```

Training and Validation Performance Difference:

Bagging: Training Score: 0.9669, Validation Score: 0.6229, Difference: 0.3440

Random forest: Training Score: 1.0000, Validation Score: 0.7082, Difference: 0.2918

GBM: Training Score: 0.7479, Validation Score: 0.7438, Difference: 0.0041

Adaboost: Training Score: 0.7015, Validation Score: 0.6985, Difference: 0.0030

dtree: Training Score: 1.0000, Validation Score: 0.6437, Difference: 0.3563

Chose the random forest, GBM and adaboost models since they were the best performing models across the board for training and validation sets and had promising results

Random Forest Model Building

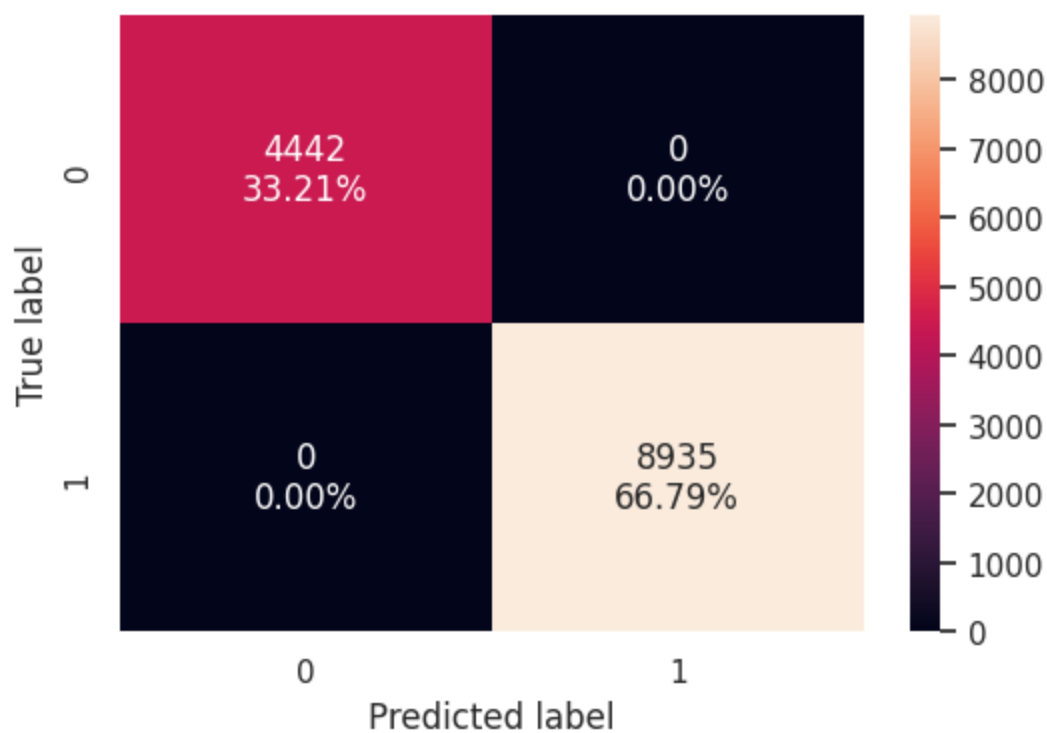
```
In [66]: rf_wt = RandomForestClassifier(class_weight='balanced', random_state=1)
rf_wt.fit(X_train,y_train)
```

Out[66]:

▼
RandomForestClassifier
(https://scikit-learn.org/1.6/module:)

RandomForestClassifier(class_weight='balanced', random_state=1)

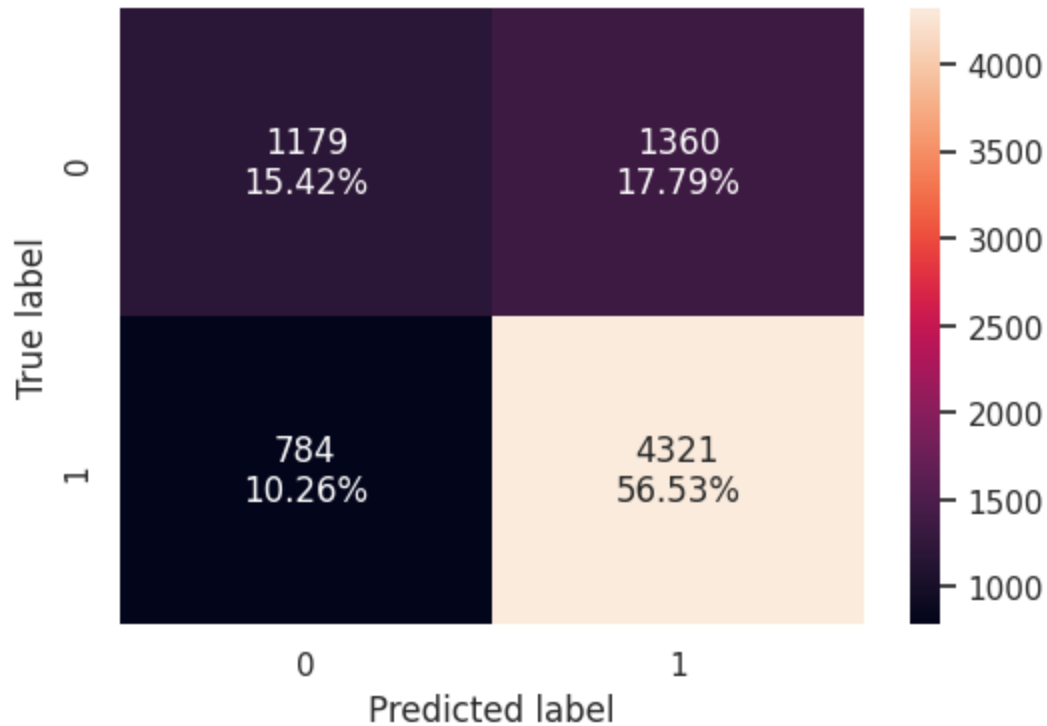
```
In [67]: confusion_matrix_sklearn(rf_wt,X_train,y_train)
```



```
In [68]: rf_wt_model_train_perf=model_performance_classification_sklearn(rf_wt, X_train,y_train)
print("Training performance \n",rf_wt_model_train_perf)
```

```
Training performance
      Accuracy  Recall  Precision    F1
0      1.000    1.000    1.000 1.000
```

```
In [69]: confusion_matrix_sklearn(rf_wt, X_test,y_test)
```



```
In [70]: rf_wt_model_test_perf=model_performance_classification_sklearn(rf_wt, X_test,y_test)
print("Testing performance \n",rf_wt_model_test_perf)

# main variable were looking at is recall which means they model is adjusting well with the dataset
```

```
Testing performance
      Accuracy  Recall  Precision   F1
0      0.720   0.846    0.761 0.801
```

Gradient Boosting Model Building

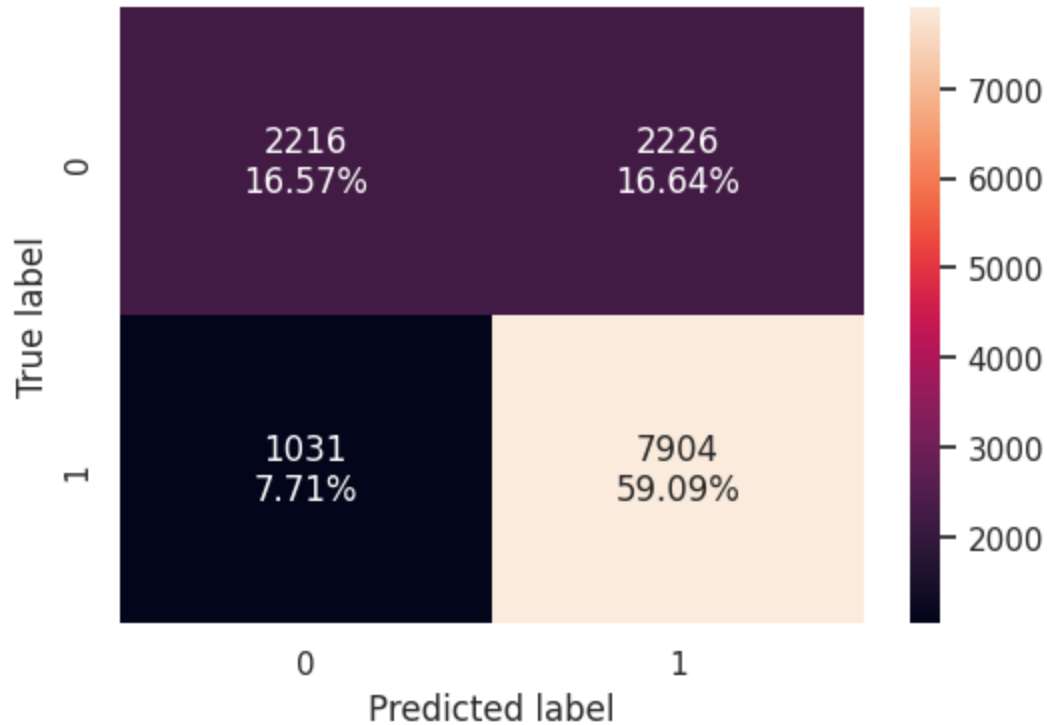
```
In [71]: gb_estimator=GradientBoostingClassifier(random_state=1)
gb_estimator.fit(X_train,y_train)
```

```
Out[71]: GradientBoostingClassifier
          (https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html)
          GradientBoostingClassifier(random_state=1)
```

```
In [72]: gb_estimator_model_train_perf = model_performance_classification_sklearn(gb_estimator, X_train, y_train)
print("Training performance \n", gb_estimator_model_train_perf)
```

```
Training performance
      Accuracy  Recall  Precision    F1
0      0.757   0.885    0.780 0.829
```

```
In [73]: confusion_matrix_sklearn(gb_estimator, X_train, y_train)
```



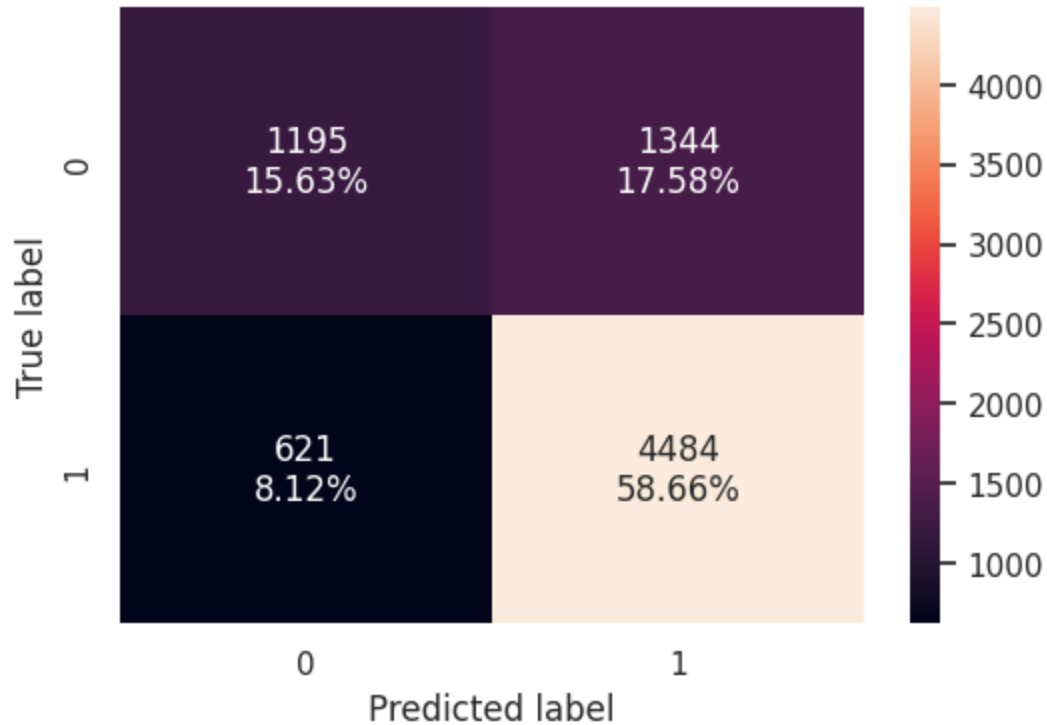
```
In [74]: gb_estimator_model_test_perf = model_performance_classification_sklearn(gb_estimator, X_test, y_test)
print("Testing performance \n", gb_estimator_model_test_perf)
```

```
Testing performance
      Accuracy  Recall  Precision    F1
0      0.743   0.878    0.769 0.820
```

```
In [75]: confusion_matrix_sklearn(gb_estimator, X_test,y_test)

# The model is performing reasonably well on both the training and test sets
# with a slight drop in accuracy and precision from training to testing, which
# is expected due to the model's generalization to unseen data

# The F1 score being similar between the train and test sets suggests that the
# model is generalizing well
# striking a good balance between false positives and false negatives.
```



AdaBoost Model Building

```
In [76]: ab_classifier = AdaBoostClassifier(random_state=1)
ab_classifier.fit(X_train, y_train)
```

```
Out[76]:
```

AdaBoostClassifier (https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.AdaBoo)

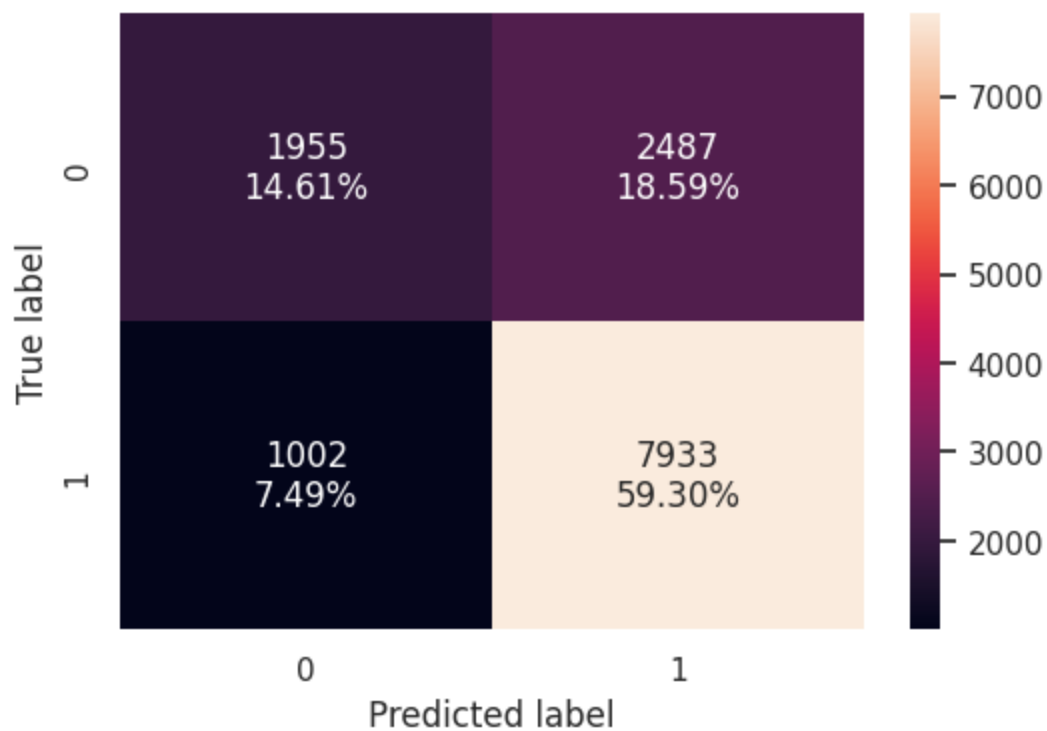
AdaBoostClassifier(random_state=1)

```
In [77]: ab_classifier_model_train_perf = model_performance_classification_sklearn(ab_c
lassifier, X_train, y_train)
print("Training performance \n", ab_classifier_model_train_perf)
```

```
Training performance
Accuracy Recall Precision F1
0 0.739 0.888 0.761 0.820
```



```
In [78]: confusion_matrix_sklearn(ab_classifier,X_train,y_train)
```

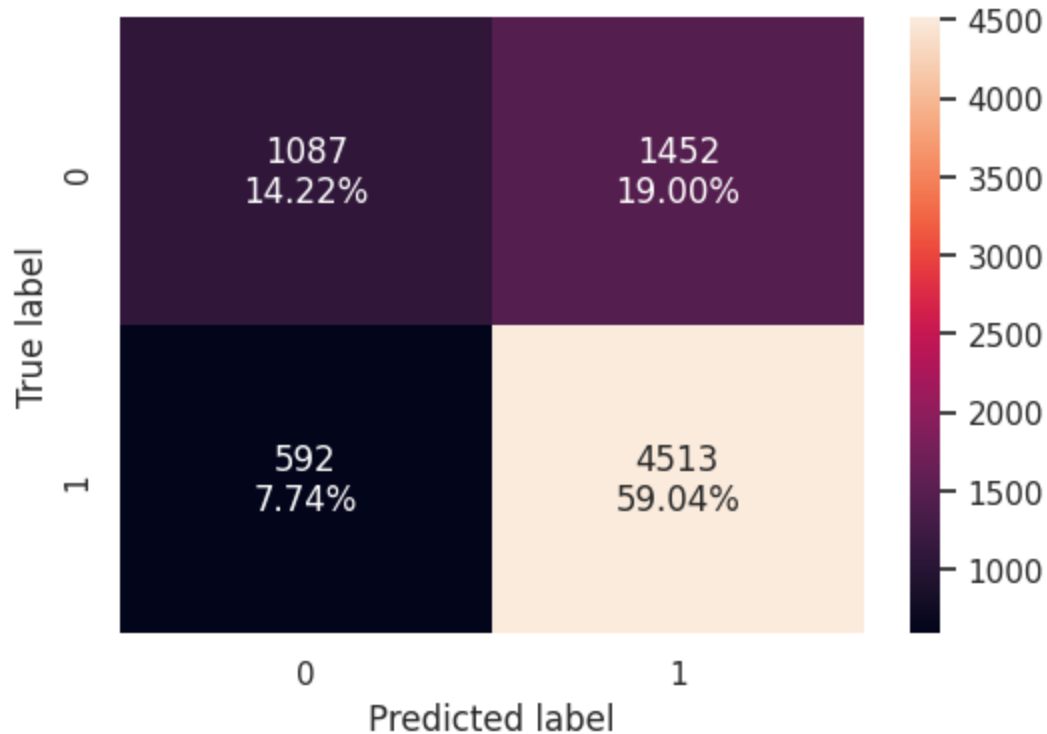


```
In [79]: ab_classifier_model_test_perf = model_performance_classification_sklearn(ab_classifier, X_test, y_test)
print("Testing performance \n", ab_classifier_model_test_perf)
```

```
Testing performance
      Accuracy  Recall  Precision    F1
0      0.733   0.884    0.757 0.815
```

```
In [80]: confusion_matrix_sklearn(ab_classifier, X_test,y_test)
```

The AdaBoost model is performing well on both the training and testing sets, with a minimal drop in performance between the two.
The similarity between training and test performance (especially recall) indicates that the model generalizes well to unseen data
The slight drop in precision from training to testing is expected, but overall, it remains effective



Hyperparameter Tuning

Tuning the GBM Model with undersampled data

```
In [81]: %%time

#Creating pipeline
Model = GradientBoostingClassifier(random_state=1)

#Parameter grid to pass in RandomSearchCV
param_grid = {
    "init": [DecisionTreeClassifier(max_depth=1, random_state=1)],
    "n_estimators": np.arange(150, 301, 50),
    "learning_rate": [0.01, 0.05, 0.1],
    "subsample": [0.7, 0.8, 0.9],
    "max_features": [0.5, 0.7],
}
scorer = metrics.make_scorer(metrics.recall_score)
```

```

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=Model,
    param_distributions=param_grid,
    n_iter=50,
    scoring=scorer,
    cv=5,
    random_state=1,
    n_jobs=-1
)

# Fitting RandomizedSearchCV
randomized_cv.fit(X_train_un, y_train_un)

print("Best parameters are {} with CV score={}".format(randomized_cv.best_params_, randomized_cv.best_score_))

```

Best parameters are {'subsample': 0.9, 'n_estimators': np.int64(200), 'max_features': 0.7, 'learning_rate': 0.01, 'init': DecisionTreeClassifier(max_depth=1, random_state=1)} with CV score=0.7615944121849634
 CPU times: user 4.96 s, sys: 594 ms, total: 5.55 s
 Wall time: 5min 36s

```

In [82]: tuned_gbm1 = GradientBoostingClassifier(
    random_state=1,
    subsample=0.9,
    n_estimators=150,
    max_features=0.5,
    learning_rate=0.01,
    init=DecisionTreeClassifier(random_state=1),
)
tuned_gbm1.fit(X_train_un, y_train_un)

```

Out[82]:

► **GradientBoostingClassifier** [?](https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html)

► **init: DecisionTreeClassifier** [?](https://scikit-learn.org/1.6/modules/generated/sklearn.tree.DecisionTreeClassifier.html)

► **DecisionTreeClassifier** [?](https://scikit-learn.org/1.6/modules/generated/sklearn.tree.DecisionTreeClassifier.html)

```

In [83]: # Checking model's performance on training set
gbm1_train = model_performance_classification_sklearn(
    tuned_gbm1, X_train_un, y_train_un
)
gbm1_train

```

Out[83]:

	Accuracy	Recall	Precision	F1
0	1.000	1.000	1.000	1.000

```
In [84]: # Checking model's performance on validation set  
gbm1_val = model_performance_classification_skllearn(tuned_gbm1, X_val, y_val)  
gbm1_val
```

Out[84]:

	Accuracy	Recall	Precision	F1
0	0.637	0.644	0.775	0.703

- The model performed perfectly on the training set but dropped significantly on the validation set, indicating overfitting
- Lower recall and F1 on the validation set suggest the model misses many positive cases.

Tuning the GBM Model with Oversampled data

```
In [85]: %%time

# Defining the model
Model = GradientBoostingClassifier(random_state=1)

# Parameter grid
param_grid = {
    "init": [DecisionTreeClassifier(max_depth=1, random_state=1)], # simpler
    "n_estimators": np.arange(150, 301, 50), # slightl
    "learning_rate": [0.01, 0.05, 0.1], # smaller
    "subsample": [0.7, 0.8, 0.9], # avoid o
    "max_features": [0.5, 0.7], # use few
}

# Scoring function (focus on recall)
scorer = metrics.make_scorer(metrics.recall_score)

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=Model,
    param_distributions=param_grid,
    n_iter=50,
    scoring=scorer,
    cv=5,
    random_state=1,
    n_jobs=-1
)

# Fitting RandomizedSearchCV on oversampled data
randomized_cv.fit(X_train_over, y_train_over)

# Printing best parameters
print("Best parameters are {} with CV score={}".format(randomized_cv.best_params_, randomized_cv.best_score_))
```

Best parameters are {'subsample': 0.7, 'n_estimators': np.int64(300), 'max_features': 0.7, 'learning_rate': 0.1, 'init': DecisionTreeClassifier(max_depth=1, random_state=1)} with CV score=0.7889199776161163
 CPU times: user 9.87 s, sys: 953 ms, total: 10.8 s
 Wall time: 10min 53s

```
In [86]: tuned_gbm2 = GradientBoostingClassifier(
    random_state=1,
    subsample=0.7,
    n_estimators=150,
    max_features=1,
    learning_rate=1,
    init=AdaBoostClassifier(random_state=1),
)
tuned_gbm2.fit(X_train_over, y_train_over)
```

Out[86]:

```

  ▸ GradientBoostingClassifier
    (https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.GradientBoos
  ▸ init: AdaBoostClassifier
    (https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.AdaBoostClassi
    ▸ AdaBoostClassifier
      (https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.AdaBoostClassi

```

```
In [87]: # Checking model's performance on training set
gbm2_train = model_performance_classification_sklearn(tuned_gbm1, X_train_over, y_train_over)
gbm2_train
```

Out[87]:

	Accuracy	Recall	Precision	F1
0	0.853	0.814	0.883	0.847

```
In [88]: # Checking model's performance on validation set
gbm2_val = model_performance_classification_sklearn(tuned_gbm1, X_val, y_val)
gbm2_val
```

Out[88]:

	Accuracy	Recall	Precision	F1
0	0.637	0.644	0.775	0.703

- The model shows good performance with high accuracy, recall, and precision, indicating a strong fit on the oversampled training data
- The validation performance drops, with a significant reduction in recall and F1 score, suggesting that the model may have overfitted to the oversampled data and struggles to generalize.

Tuning RandomForestClassifier model with undersampled data

```
In [89]: %%time

# Defining the model
Model = RandomForestClassifier(random_state=1)

# Parameter grid
param_grid = {
    "n_estimators": np.arange(150, 301, 50),      # Number of trees
    "max_depth": [5, 10, 15, None],              # Depth of each tree
    "min_samples_split": [2, 5, 10],             # Min samples to split a node
    "min_samples_leaf": [1, 2, 4],               # Min samples at a leaf node
    "max_features": [0.5, 0.7, 1],               # Number of features to consider at split
    "class_weight": [None, 'balanced']           # Try class balancing
}

# Scoring function (focus on recall)
scorer = metrics.make_scorer(metrics.recall_score)

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=Model,
    param_distributions=param_grid,
    n_iter=50,
    scoring=scorer,
    cv=5,
    random_state=1,
    n_jobs=-1
)

# Fitting RandomizedSearchCV on undersampled data
randomized_cv.fit(X_train_un, y_train_un)

# Printing best parameters
print("Best parameters are {} with CV score={}".format(randomized_cv.best_params_, randomized_cv.best_score_))
```

```
Best parameters are {'n_estimators': np.int64(250), 'min_samples_split': 2,
'min_samples_leaf': 4, 'max_features': 0.7, 'max_depth': 10, 'class_weight':
'balanced'} with CV score=0.7485402162567516
CPU times: user 8.72 s, sys: 690 ms, total: 9.41 s
Wall time: 9min 4s
```

```
In [90]: # Train the final tuned Random Forest model
tuned_rf1 = RandomForestClassifier(
    random_state=1,
    n_estimators=150,
    max_depth=10,
    min_samples_split=2,
    min_samples_leaf=1,
    max_features=0.7,
    class_weight='balanced'
)

tuned_rf1.fit(X_train_un, y_train_un)
```

```
Out[90]: RandomForestClassifier
RandomForestClassifier(class_weight='balanced', max_depth=10, max_features=
0.7,
                        n_estimators=150, random_state=1)
```

```
In [91]: # Checking model's performance on training set
rf1_train = model_performance_classification_sklearn(
    tuned_rf1, X_train_un, y_train_un
)
rf1_train
```

```
Out[91]:
```

	Accuracy	Recall	Precision	F1
0	0.800	0.845	0.776	0.809

```
In [92]: # Checking model's performance on validation set
rf1_val = model_performance_classification_sklearn(tuned_rf1, X_val, y_val)
rf1_val
```

```
Out[92]:
```

	Accuracy	Recall	Precision	F1
0	0.724	0.758	0.815	0.786

- The model demonstrates balanced performance with decent accuracy, recall, precision, and F1 score, indicating good learning on the training data
- The performance remains consistent with the training set, suggesting that the model is generalizing well and not overfitting.

Tuning RandomForestClassifier model with oversampled data

```
In [93]: %%time

# Defining the model
Model = RandomForestClassifier(random_state=1)

# Parameter grid
param_grid = {
    "n_estimators": np.arange(150, 301, 50),      # Number of trees
    "max_depth": [5, 10, 15, None],              # Depth of each tree
    "min_samples_split": [2, 5, 10],              # Minimum samples to split a node
    "min_samples_leaf": [1, 2, 4],                # Minimum samples at a leaf node
    "max_features": [0.5, 0.7, 1],                # Features considered at each split
    "class_weight": [None, "balanced"]            # Try class balancing
}

# Scoring function (focus on recall)
scorer = metrics.make_scorer(metrics.recall_score)

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=Model,
    param_distributions=param_grid,
    n_iter=50,
    scoring=scorer,
    cv=5,
    random_state=1,
    n_jobs=-1
)

# Fitting RandomizedSearchCV on oversampled data
randomized_cv.fit(X_train_over, y_train_over)

# Printing best parameters
print("Best parameters are {} with CV score={}".format(randomized_cv.best_params_, randomized_cv.best_score_))
```

```
Best parameters are {'n_estimators': np.int64(250), 'min_samples_split': 10,
'min_samples_leaf': 1, 'max_features': 1, 'max_depth': 15, 'class_weight': None}
with CV score=0.7985450475657527
CPU times: user 11.3 s, sys: 1.46 s, total: 12.8 s
Wall time: 17min 36s
```

```
In [94]: # Train the final tuned Random Forest model on oversampled data
tuned_rf2 = RandomForestClassifier(
    random_state=1,
    n_estimators=150,
    max_depth=10,
    min_samples_split=2,
    min_samples_leaf=1,
    max_features=0.7,
    class_weight='balanced'
)

tuned_rf2.fit(X_train_over, y_train_over)
```

```
Out[94]: RandomForestClassifier
RandomForestClassifier(class_weight='balanced', max_depth=10, max_features=
0.7,
                        n_estimators=150, random_state=1)
```

```
In [95]: # Checking model's performance on training set
rf2_train = model_performance_classification_sklearn(
    tuned_rf2, X_train_over, y_train_over
)
rf2_train
```

```
Out[95]:
```

	Accuracy	Recall	Precision	F1
0	0.823	0.841	0.812	0.826

```
In [96]: # Checking model's performance on validation set
rf2_val = model_performance_classification_sklearn(tuned_rf2, X_val, y_val)
rf2_val
```

```
Out[96]:
```

	Accuracy	Recall	Precision	F1
0	0.739	0.811	0.801	0.806

- The model performs well with high accuracy, recall, and F1 score, suggesting it has effectively learned from the oversampled data.
- While the validation performance is slightly lower than training, the recall and F1 score are still strong, showing that the model is maintaining good generalization.

```
In [ ]:
```

Tuning AdaBoost model with undersampled data

```
In [97]: %%time

# defining model
Model = AdaBoostClassifier(random_state=1)

# Parameter grid to pass in RandomSearchCV
param_grid = {
    "n_estimators": np.arange(10, 40, 10),
    "learning_rate": [0.1, 0.01, 0.2, 0.05, 1],
    "estimator": [
        DecisionTreeClassifier(max_depth=1, random_state=1),
        DecisionTreeClassifier(max_depth=2, random_state=1),
        DecisionTreeClassifier(max_depth=3, random_state=1),
    ],
}

# Type of scoring used to compare parameter combinations
scorer = metrics.make_scorer(metrics.recall_score)

#Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(estimator=Model, param_distributions=param_
grid, n_jobs = -1, n_iter=50, scoring=scorer, cv=5, random_state=1)

#Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_un, y_train_un)

print("Best parameters are {} with CV score={}" .format(randomized_cv.best_pa
rams_,randomized_cv.best_score_))
```

Best parameters are {'n_estimators': np.int64(30), 'learning_rate': 0.05, 'es
timator': DecisionTreeClassifier(max_depth=1, random_state=1)} with CV score=
0.8036930349922476:

CPU times: user 985 ms, sys: 113 ms, total: 1.1 s

Wall time: 37 s

```
In [98]: tuned_adb = AdaBoostClassifier(
    random_state=1,
    n_estimators=20,
    learning_rate=0.1,
    estimator=DecisionTreeClassifier(max_depth=2, random_state=1),
)
tuned_adb.fit(X_train_un, y_train_un)
```

Out[98]:

AdaBoostClassifier

estimator:

DecisionTreeClassifier

DecisionTreeClassifier

(<https://scikit-learn.org/1.6/modules/generated/sklearn.ensemble.AdaBo>

(<https://scikit-learn.org/1.6/modules/generated/sklearn.tree.DecisionTreeC>

```
In [99]: # Checking model's performance on training set
adb_train = model_performance_classification_sklearn(tuned_adb, X_train_un, y_train_un)
adb_train
```

Out[99]:

	Accuracy	Recall	Precision	F1
0	0.684	0.782	0.654	0.712

```
In [100]: # Checking model's performance on validation set
adb_val = model_performance_classification_sklearn(tuned_adb, X_val, y_val)
adb_val
```

Out[100]:

	Accuracy	Recall	Precision	F1
0	0.721	0.781	0.796	0.789

- The model shows moderate accuracy and recall with a balanced F1 score, indicating it learns the data but might have room for improvement in precision
- The validation performance improves slightly with a solid recall and precision score, suggesting good generalization while maintaining a balanced F1 score.

Model Performances

```
In [101]: # training performance comparison

models_train_comp_df = pd.concat(
    [
        gbm1_train.T,
        gbm2_train.T,
        rf1_train.T,
        rf2_train.T,
        adb_train.T,
    ],
    axis=1,
)
models_train_comp_df.columns = [
    "GBM trained with Undersampled data",
    "GBM trained with Oversampled data",
    "Random Forest trained with Undersampled data",
    "Random Forest trained with Oversampled data",
    "AdaBoost Forest trained with Undersampled data",
]
print("Training performance comparison:")
models_train_comp_df
```

Training performance comparison:

Out[101]:

	GBM trained with Undersampled data	GBM trained with Oversampled data	Random Forest trained with Undersampled data	Random Forest trained with Oversampled data	AdaBoost Forest trained with Undersampled data
Accuracy	1.000	0.853	0.800	0.823	0.684
Recall	1.000	0.814	0.845	0.841	0.782
Precision	1.000	0.883	0.776	0.812	0.654
F1	1.000	0.847	0.809	0.826	0.712

-GBM trained with Undersampled data: This model achieved perfect performance across all metrics, but it likely overfits the training data, which might not generalize well.

-GBM trained with Oversampled data: The model showed solid performance with high accuracy and balanced recall, precision, and F1 score, suggesting better generalization on the oversampled dataset.

-Random Forest trained with Undersampled data: The model performed well with high recall, though accuracy and precision were slightly lower, indicating it might be biased toward detecting positive cases.

-Random Forest trained with Oversampled data: With a slight improvement in accuracy and F1 score, this model demonstrated better performance compared to the undersampled version.

-AdaBoost trained with Undersampled data: This model had the lowest performance across all metrics, particularly in recall and precision, indicating that it struggles to handle the undersampled data effectively.

```
In [102]: # Validation performance comparison

models_train_comp_df = pd.concat(
    [ gbm1_val.T,
      gbm2_val.T,
      rf1_val.T,
      rf2_val.T,
      adb_val.T,
    ]
    , axis=1,
)
models_train_comp_df.columns = [
    "GBM trained with Undersampled data",
    "GBM trained with Oversampled data",
    "Random Forest trained with Undersampled data",
    "Random Forest trained with Oversampled data",
    "AdaBoost Forest trained with Undersampled data",
]
print("Validation performance comparison:")
models_train_comp_df
```

Validation performance comparison:

Out[102]:

	GBM trained with Undersampled data	GBM trained with Oversampled data	Random Forest trained with Undersampled data	Random Forest trained with Oversampled data	AdaBoost Forest trained with Undersampled data
Accuracy	0.637	0.637	0.724	0.739	0.721
Recall	0.644	0.644	0.758	0.811	0.781
Precision	0.775	0.775	0.815	0.801	0.796
F1	0.703	0.703	0.786	0.806	0.789

-GBM trained with Undersampled data: The validation results show moderate performance, with balanced recall and precision, but the model's accuracy and F1 score could be improved.

-GBM trained with Oversampled data: Similar to the undersampled version, this model shows balanced recall and precision but could benefit from higher accuracy and F1 score.

-Random Forest trained with Undersampled data: This model performed better on the validation set with a noticeable improvement in recall and precision, leading to a higher F1 score than the GBM models.

-Random Forest trained with Oversampled data: The model showed a slight increase in recall and F1 score, suggesting that oversampling has improved its generalization and classification performance.

-AdaBoost trained with Undersampled data: With strong recall and precision, this model performed well in terms of F1 score, though there is still room for improvement in accuracy.

```
In [103]: # Let's check the performance on test set
rf_test = model_performance_classification_sklern(tuned_rf1, X_test, y_test)
rf_test
```

Out[103]:

	Accuracy	Recall	Precision	F1
0	0.708	0.742	0.805	0.772

```
In [104]: # Let's check the performance on test set
rf_test2 = model_performance_classification_sklern(tuned_rf2, X_test, y_test)
rf_test2

#winner for best model with highest recall score, others were close but random
forest did good several times and ended up being #1
```

Out[104]:

	Accuracy	Recall	Precision	F1
0	0.716	0.791	0.786	0.788

```
In [105]: # Let's check the performance on test set
gbm_test = model_performance_classification_sklern(tuned_gbm1, X_test, y_test)
gbm_test
```

Out[105]:

	Accuracy	Recall	Precision	F1
0	0.620	0.629	0.760	0.688

```
In [106]: # Let's check the performance on test set
gbm_test2 = model_performance_classification_sklern(tuned_gbm2, X_test, y_test)
gbm_test2
```

Out[106]:

	Accuracy	Recall	Precision	F1
0	0.702	0.774	0.779	0.776

```
In [107]: # Let's check the performance on test set
ada_test = model_performance_classification_sklern(tuned_adb, X_test, y_test)
ada_test
```

Out[107]:

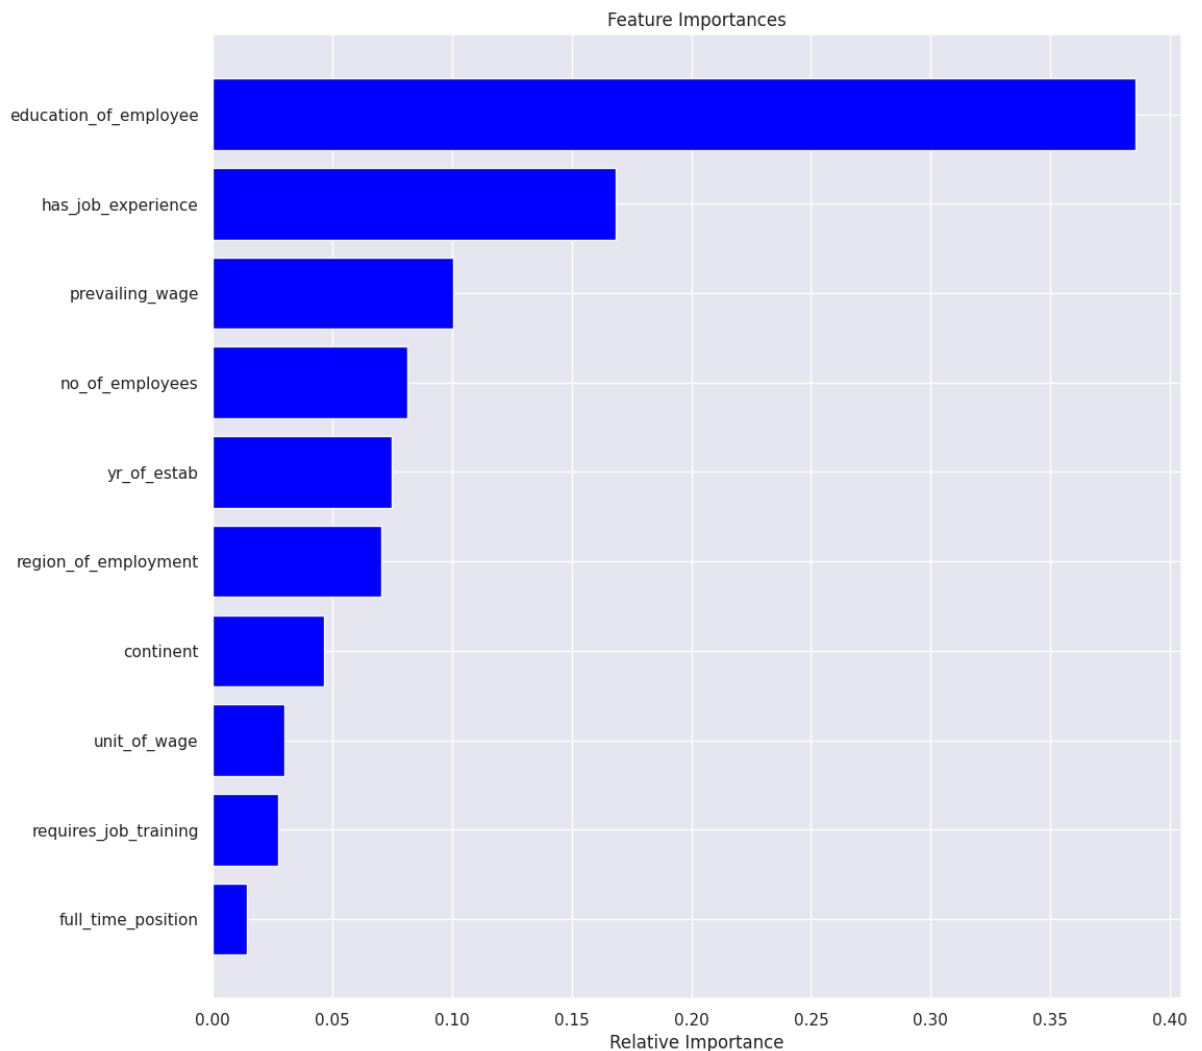
	Accuracy	Recall	Precision	F1
0	0.714	0.780	0.789	0.784

- The best performing model appears to be rf_test2 with the highest recall score (0.791)
- AdaBoost is a close second with an F1 score of 0.780
- The first gradient boosting model (gbm_test) underperformed significantly but improved dramatically after tuning
- All models except the first gbm_test have fairly balanced precision and recall

Feature Importance

```
In [108]: feature_names = X_train.columns
importances = tuned_rf2.feature_importances_
indices = np.argsort(importances)

plt.figure(figsize=(12, 12))
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="blue", align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```



Actionable Insights and Recommendations

Key Insights & Suggestions

- Best Model: The Random Forest model (rf_test2) did the best job overall, with a solid F1 score of 0.788. It balanced precision (0.786) and recall (0.791) well, making it the most reliable for predicting visa approvals.
- Close Second: The AdaBoost model also did well, with an F1 score of 0.784 and slightly better precision (0.789) than the Random Forest. It's a good backup option.
- Class Imbalance Matters: The model worked a lot better after adjusting for class imbalance (using SMOTE), which shows that handling uneven visa application data is key for good predictions.

Suggestions for Improving the Process:

- Two Stage Screening:
- Use the Random Forest model as the main tool to screen applications.
- For tricky cases, double check with the AdaBoost model to avoid wrong denials. This way, you get better accuracy and fairness in visa decisions.

Focus on High-Risk Applications:

- Applications that are likely to be denied should get a second look by a person.
- Fast track those with a high chance of approval to save time. This helps use resources more efficiently and reduce backlogs.

Create a Pre-Screening Tool:

- Build an online tool where people can check their approval chances before applying. This cuts down on applications that are unlikely to be approved and eases the workload.

Applicant Profile Suggestions: Based on the model's findings, applicants with these characteristics are more likely to get approved:

Wage Compliance:

- Make sure the wage offered meets or exceeds the industry standard. Low wages are linked to higher denials.

High-Demand Areas:

- Jobs in regions with labor shortages tend to get approved more. Focus on these areas when applying.

Education and Experience:

- Applicants whose education and experience match the job better have a higher chance of approval.

Full-Time Jobs:

- Full-time positions are more likely to get approved compared to part-time or contract jobs.

How to Roll This Out:

Start Small:

- Begin by using the Random Forest model on a small batch of applications.
- Compare how well it works against the current manual process.
- If it improves efficiency, expand it.

Keep Improving the Model:

- Regularly retrain the model with new data.
- Review the model every few months to keep its predictions sharp.
- Try A/B testing different models to see which one works best

Power Ahead
